

Machine Learning with Python



Agenda

- What is Machine Learning
- Supervised Learning
 - Regression
 - Classification
- Unsupervised Learning
 - Clustering
 - Dimension Reduction
- Model Selection & Evaluation
 - Cross Validation
 - Hyperparameter Tuning
- Intro to Deep Learning




Agenda

- What is Machine Learning
- Supervised Learning
 - Regression
 - Classification
- Unsupervised Learning
 - Clustering
 - Dimension Reduction
- Model Selection & Evaluation
 - Cross Validation
 - Hyperparameter Tuning
- Intro to Deep Learning

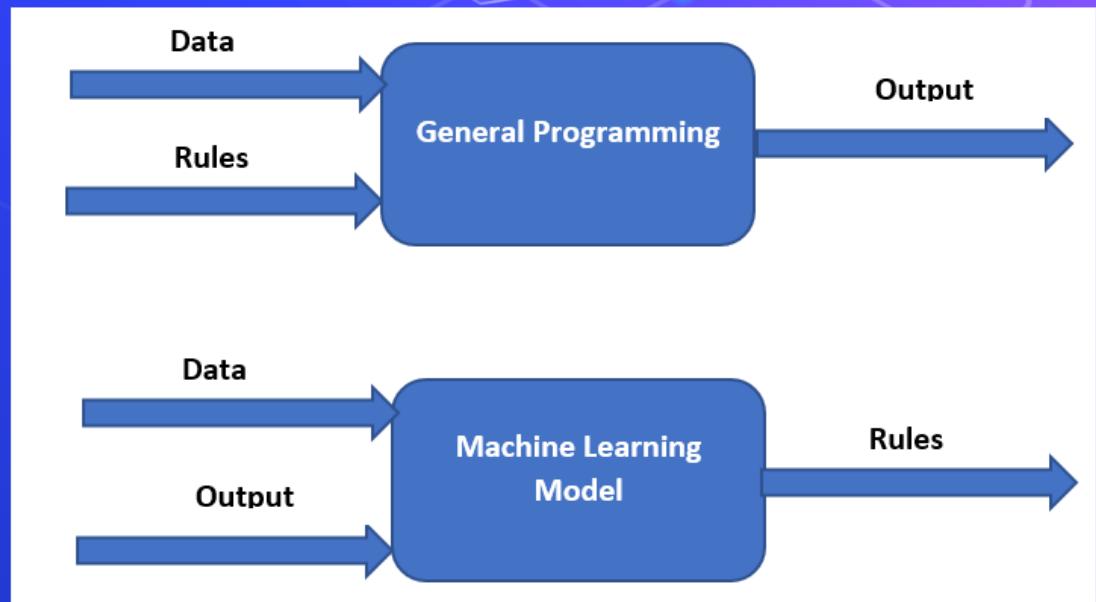
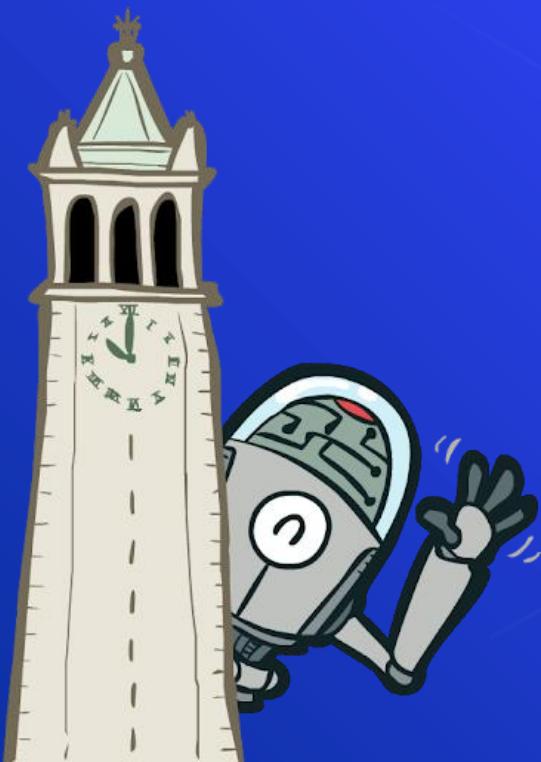


What is Machine Learning

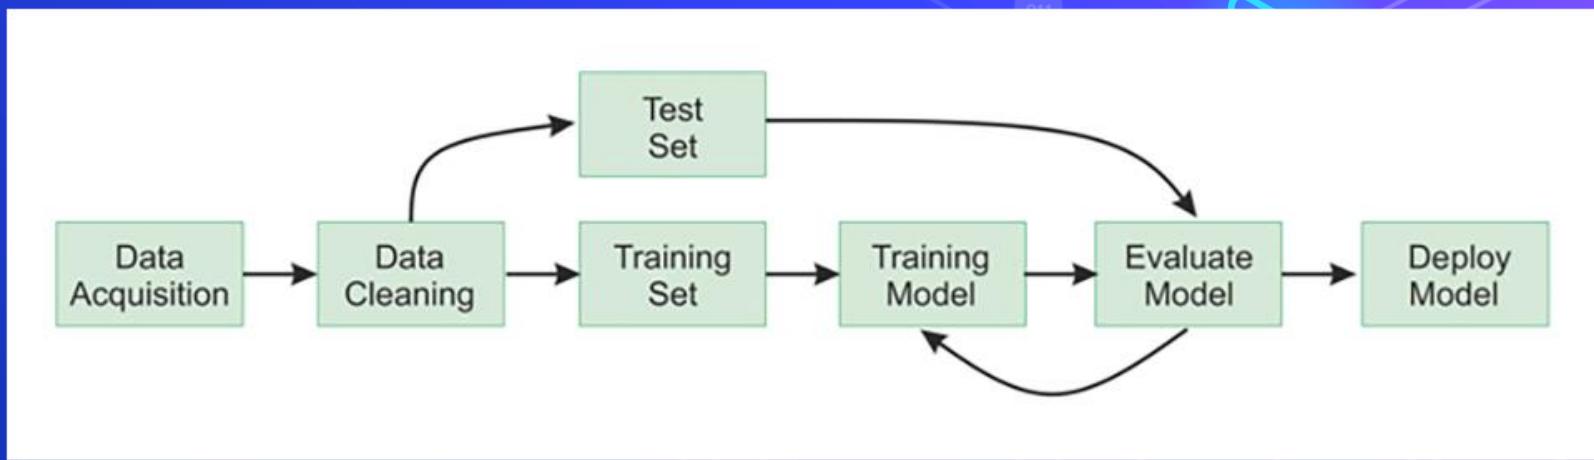
Machine learning involves computers discovering how they can perform tasks without being explicitly programmed to do so. It involves computers learning from data provided so that they carry out certain tasks.



What is Machine Learning



What is Machine Learning



What is Machine Learning

Classical Machine Learning

Task Driven

Supervised Learning

(Pre Categorized Data)



Classification

(Divide the socks by Color)

Eg. Identity Fraud Detection



Regression

(Divide the Ties by Length)

Eg. Market Forecasting

Data Driven

Unsupervised Learning

(Unlabelled Data)



Clustering

(Divide by Similarity)

Eg. Targeted Marketing



Association

(Identify Sequences)

Eg. Customer Recommendation



Dimensionality Reduction

(Wider Dependencies)

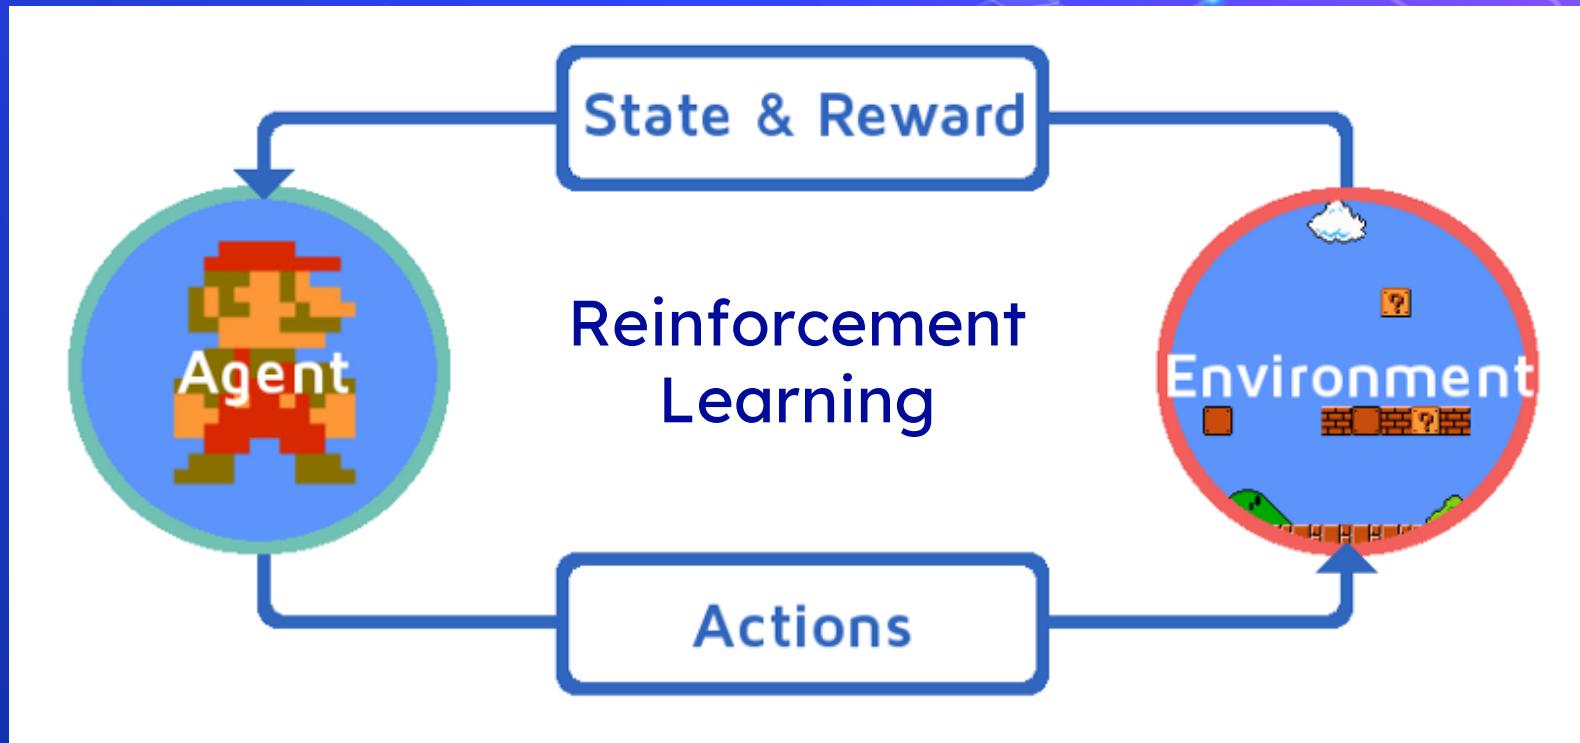
Eg. Big Data Visualization

Obj: Predictions & Predictive Models

Pattern/ Structure Recognition



What is Machine Learning



What is Machine Learning

- Supervised learning : Task Driven (Classification, Regression)
- Unsupervised learning : Data Driven (Clustering)
- Reinforcement learning :
 - Close to human learning.
 - Algorithm learns a policy of how to act in a given environment.
 - Every action has some impact in the environment, and the environment provides rewards that guides the learning algorithm.

Agenda

- What is Machine Learning
- **Supervised Learning**
 - Regression
 - Classification
- Unsupervised Learning
 - Clustering
 - Dimension Reduction
- Model Selection & Evaluation
 - Cross Validation
 - Hyperparameter Tuning
- Intro to Deep Learning

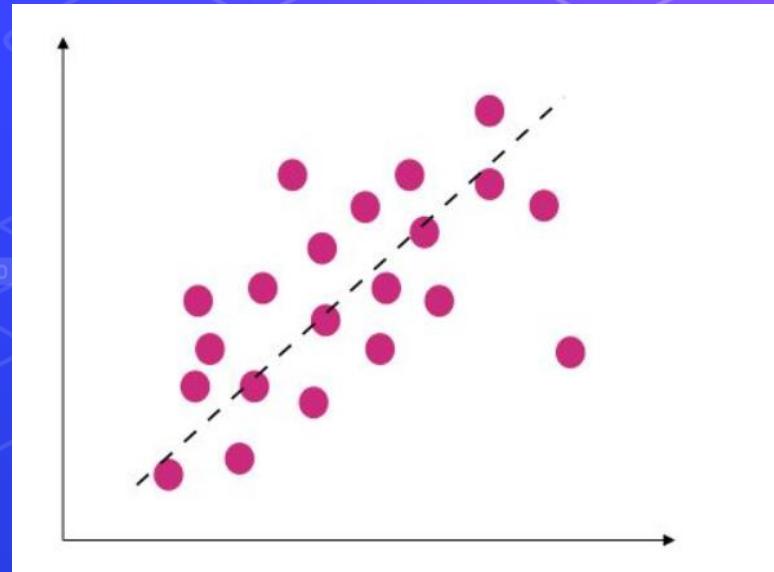


Regression

Regression quantifies the relationship between one or more predictor variable(s) and one outcome variable.

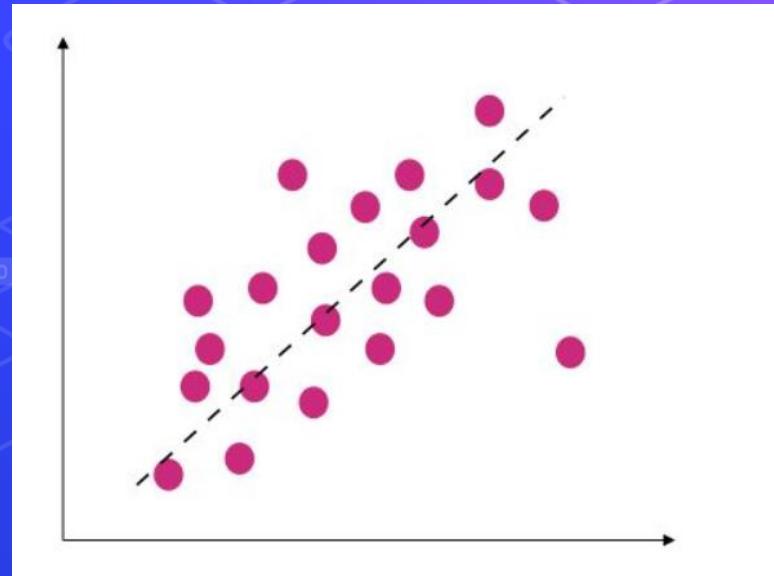
For example,
it can be used to quantify the relative impacts of age, gender,
and diet (the predictors) on weight (the output or dependent).

- House prices based on size, locations,...
- Salary prediction based on experience
- Stock Market prediction
- Weather prediction
- ...




Regression

- Simple & Multiple Linear Regression
- Other Regression Algorithms
- Evaluating Model Performance




Regression

- ❖ Simple & Multiple Linear Regression
- ❖ Other Regression Algorithms
- ❖ Evaluating Model Performance



Simple Linear Regression (Equation of a straight line)

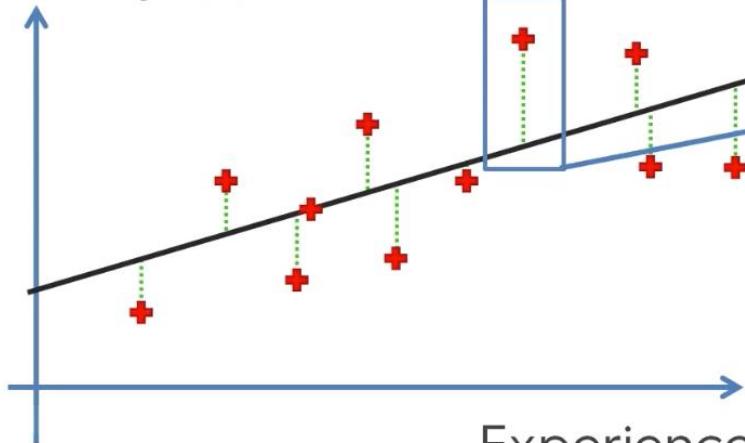
It also called ordinary least squares (OLS)



Simple Linear Regression (Calculate Cost Function)

Simple Linear Regression:

Salary (\$)



y_i

\hat{y}_i

Hypothesis:
$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

Parameters: θ_0, θ_1

Cost Function:
$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Goal:
$$\underset{\theta_0, \theta_1}{\text{minimize}} J(\theta_0, \theta_1)$$

Simple Linear Regression (Calculate Cost Function)

The goal is to minimize the sum of squares of residuals ($y - y'$) is as small as possible.

Hypothesis: $h_{\theta}(x) = \theta_0 + \theta_1 x$

Parameters: θ_0, θ_1

Cost Function: $J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$

Goal: $\underset{\theta_0, \theta_1}{\text{minimize}} J(\theta_0, \theta_1)$



Simple Linear Regression (Example)

Theta0 = 5 , theta 1 = 2

Equation $h(x) = 5 + 2x$


X	Y	$h(x)$	$h(x) - y$	$(h(x) - y)^2$
1	7	7	0	0
2	8	9	1	1
2	7	9	2	4
3	9	11	2	4
4	11	13	2	4
5	10	15	5	25
5	12	15	3	9

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)})^2$$

$$J = 1 / 14 (0+1+4+4+4+25+9)$$

$$J = 47/14 = 3.3$$

Simple Linear Regression (Minimize cost using Gradient Descent)



repeat until convergence:

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$$

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_\theta(x_i) - y_i)$$

$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m ((h_\theta(x_i) - y_i)x_i)$$

Simple Linear Regression (Minimize cost using Gradient Descent)

Cost Function

$$J(\Theta_0, \Theta_1) = \frac{1}{2m} \sum_{i=1}^m [h_\Theta(x_i) - y_i]^2$$

↑
Predicted Value ↑
True Value

Gradient Descent

$$\Theta_j = \Theta_j - \alpha \frac{\partial}{\partial \Theta_j} J(\Theta_0, \Theta_1)$$

↑
Learning Rate

NOW,

$$\begin{aligned}\frac{\partial}{\partial \Theta} J_\Theta &= \frac{\partial}{\partial \Theta} \frac{1}{2m} \sum_{i=1}^m [h_\Theta(x_i) - y_i]^2 \\&= \frac{1}{m} \sum_{i=1}^m (h_\Theta(x_i) - y) \frac{\partial}{\partial \Theta_j} (\Theta x_i - y) \\&= \frac{1}{m} (h_\Theta(x_i) - y) x_i\end{aligned}$$


Therefore,

$$\Theta_j := \Theta_j - \frac{\alpha}{m} \sum_{i=1}^m [(h_\Theta(x_i) - y) x_i]$$




Simple Linear Regression (Alpha constant)

Big learning rate



Small learning rate



Multiple Linear Regression (Equation)

Simple
Linear
Regression

$$y = b_0 + b_1 * x_1$$

Multiple
Linear
Regression

Dependent variable (DV) Independent variables (IVs)

$$y = b_0 + b_1 * x_1 + b_2 * x_2 + \dots + b_n * x_n$$

Constant Coefficients

```
graph TD; DV[Dependent variable DV] --> y; IVs[Independent variables IVs] --> x1; IVs --> x2; IVs --> xn; Constant[Constant] --> b0; Coefficients[Coefficients] --> b1x1; Coefficients --> b2x2; Coefficients --> bnxn;
```

Multiple Linear Regression (Update theta values)

repeat until convergence: {

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) \cdot x_0^{(i)}$$

$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) \cdot x_1^{(i)}$$

$$\theta_2 := \theta_2 - \alpha \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) \cdot x_2^{(i)}$$

...

}

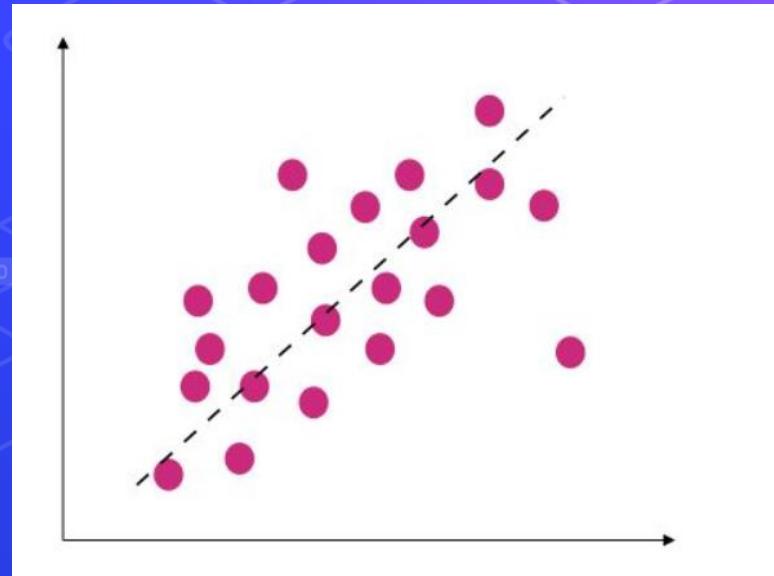
Linear Regression (Code)

```
 1 from sklearn.linear_model import LinearRegression  
 2  
 3 # make model object  
 4 model = LinearRegression()  
 5  
 6 # train  
 7 model.fit(x_train, y_train)  
 8  
 9 # test  
10 model.predict(x_test)  
11  
12 # calculate R2 score on training data  
13 model.score(x_train, y_train)  
14  
15 # calculate R2 score on testing data  
16 model.score(x_test, y_test)  
17  
18 # get model parameters  
19 model.coef_  
20 model.intercept_
```



Regression

- Simple & Multiple Linear Regression
- Other Regression Algorithms
- Evaluating Model Performance


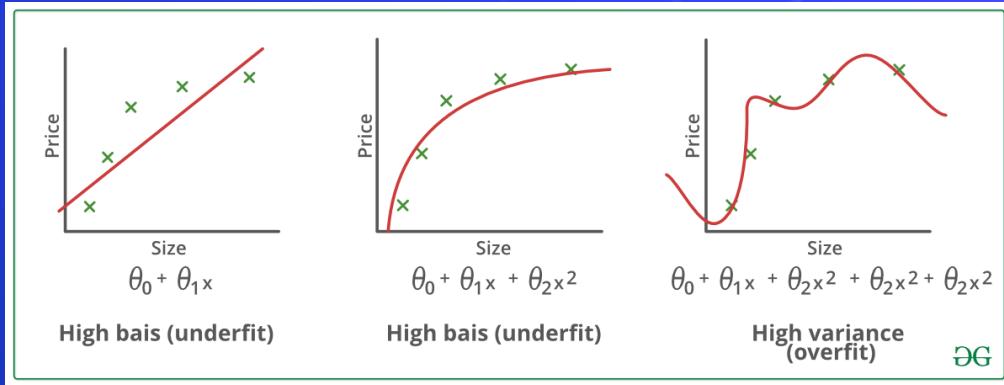


Other Regression Algorithms

- K-Nearest Neighbors (KNN)
- SVM
- Decision Trees
- Random Forests
- XGBoost



Underfitting & Overfitting

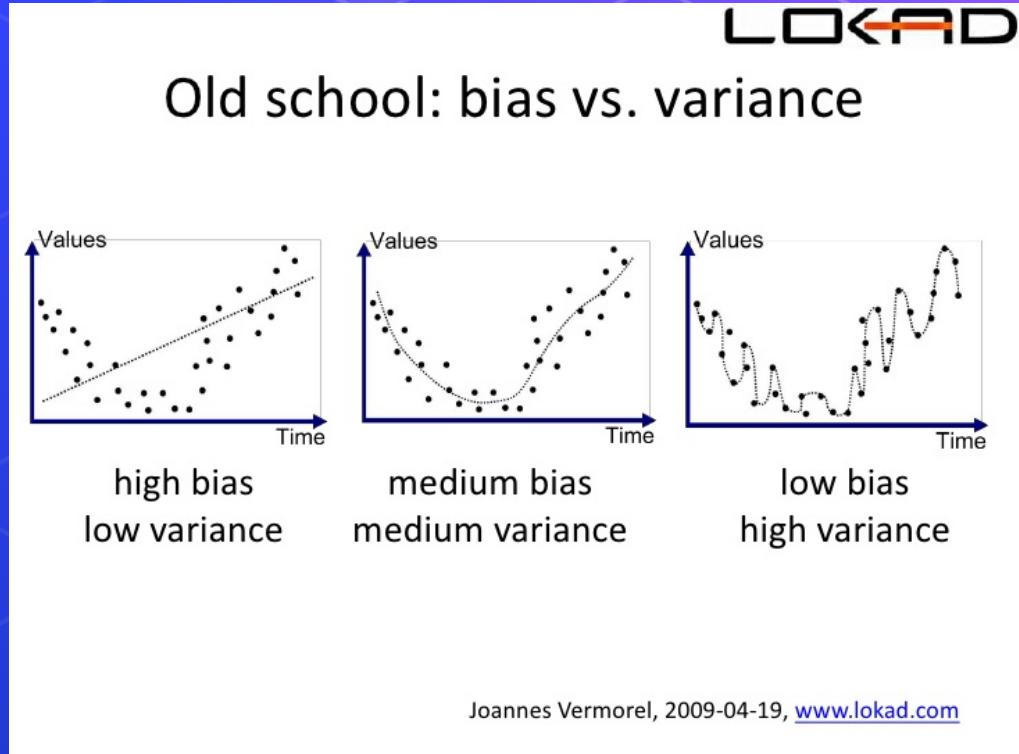


Underfitting & Overfitting

The more features you add the accuracy improves but this can lead to **Overfitting** that means the model memorizes the data too much and not applying **Generalization**.

So we should know about the trade-off between **Bias** and **Variance**.

- A model with a high bias error underfits data and makes very simplistic assumptions on it.
- A model with a high variance error overfits the data and learns too much from it.
- A good model is where both Bias and Variance errors are balanced



Underfitting & Overfitting

Underfitting (High Bias, Low Variance)

Model is so simple and weak to understand and represent data.

High Train Error + High Test Error

Overfitting (Low Bias, High Variance)

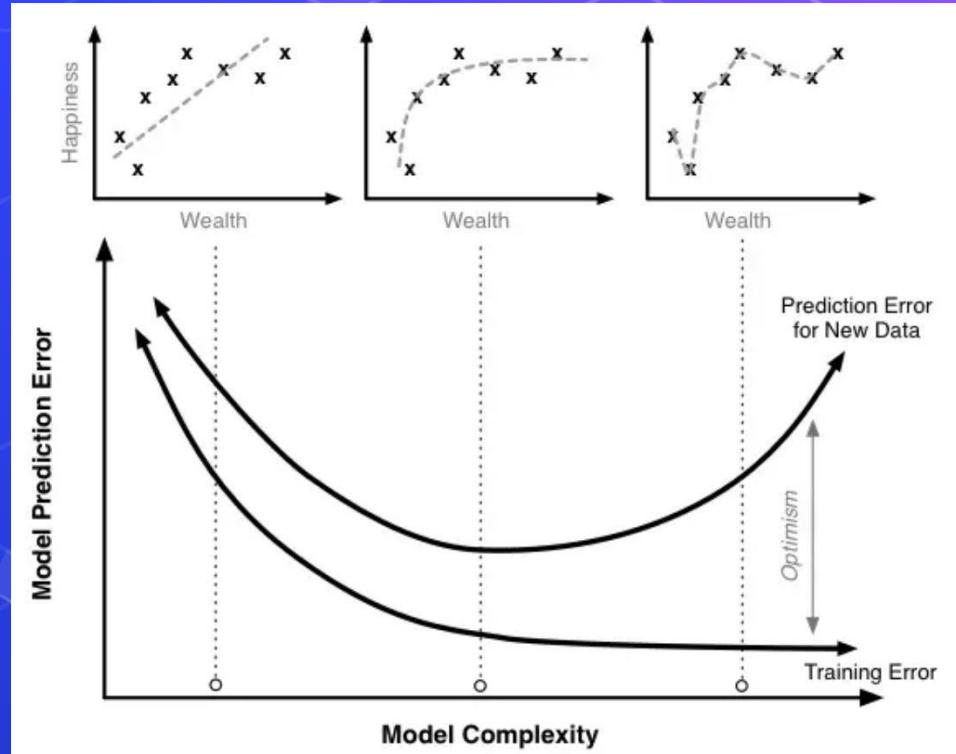
Model is so complex so it memorizes training data and can't apply generalization.

Low Train Error + High Test Error

Optimal (Low Bias, Low Variance)

Model is not simple nor complex and apply generalization well.

Low Train Error + Low Test Error



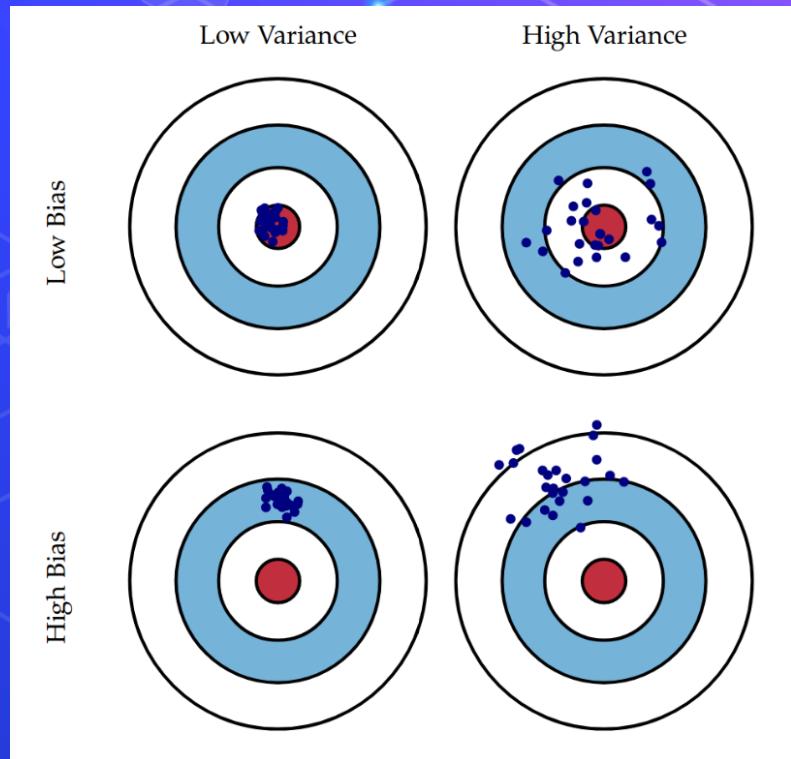
Underfitting & Overfitting

Techniques to reduce underfitting

- Increase model complexity.
- Increase number of features.
- Performing feature engineering.
- Remove noise from the data.

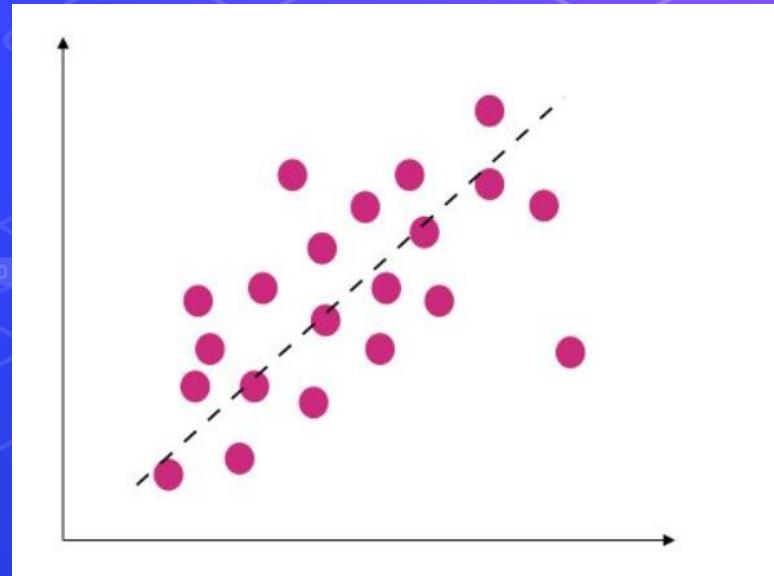
Techniques to reduce overfitting

- Increase training data.
- Remove correlated features.
- Use simpler model.
- Regularization (add penalty bias).



Regression

- Simple & Multiple Linear Regression
- Other Regression Algorithms
- Evaluating Model Performance



Evaluating Model Performance (R²)

Hrs Studied (X)	Marks (Y)
0	40
2	52
3	53
4	55
4	56
5	72
6	71
6	88
7	56
7	74
8	89
9	67
9	89
5.38	66.31
Mean	

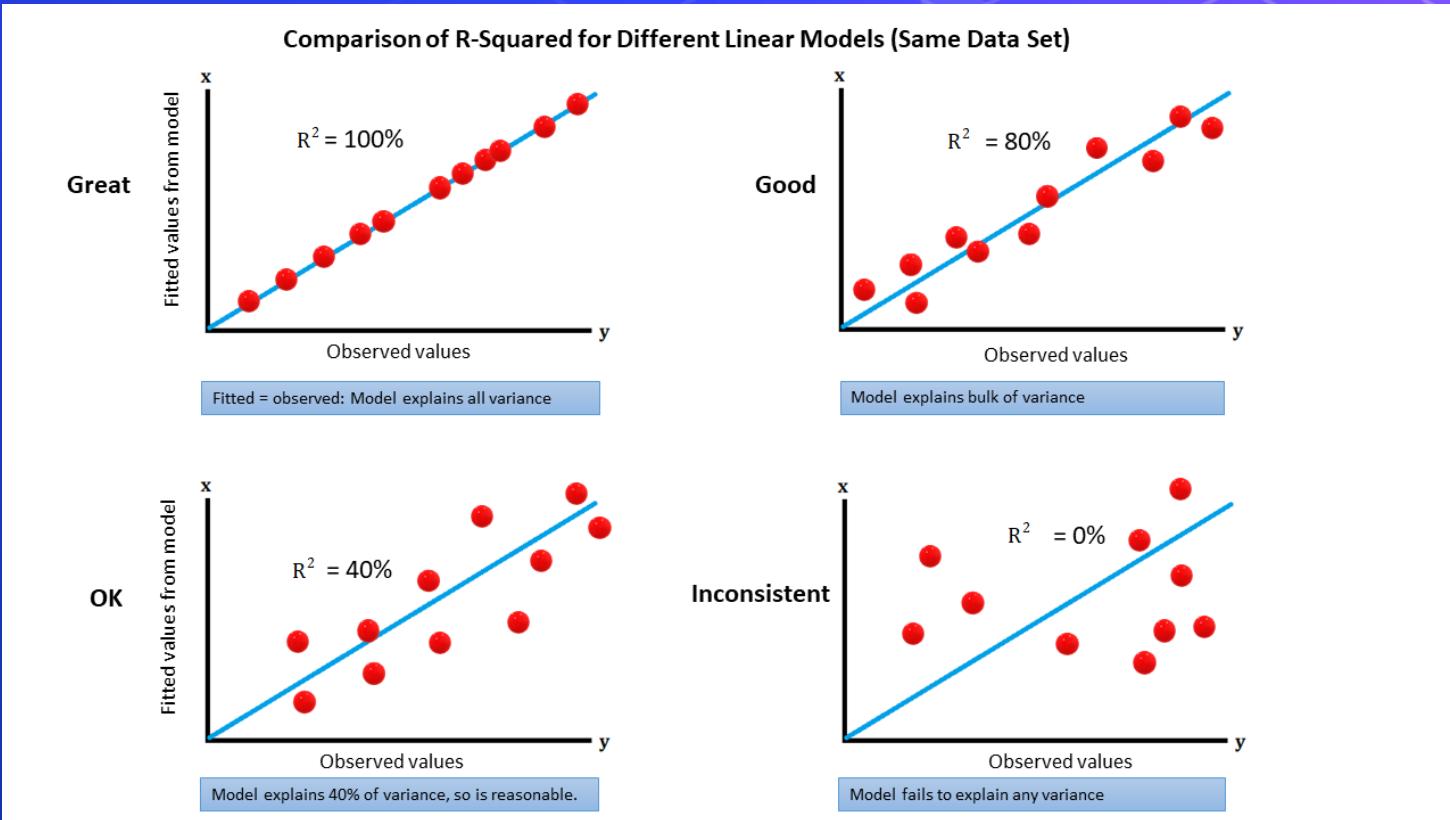
$$\begin{aligned} R^2 &= \text{SSR/SST} \\ &= 1844.12/3028.77 \\ &= 0.60886 \end{aligned}$$

Higher the value → Variation in Y is explained by variation in X.

$$\hat{y} \rightarrow y \quad \text{SSR} \rightarrow \text{SST} \quad R^2 \approx 1$$

$(Y - \bar{Y})^2$	$(\hat{Y} - \bar{Y})^2$
692.22	600.74
204.78	237.47
177.16	117.94
127.92	39.82
106.30	39.82
32.38	3.10
22.00	7.78
470.46	7.78
106.30	53.88
59.14	53.88
514.84	141.37
0.48	270.27
514.84	270.27
3028.77	1844.12
SST	SSR

Evaluating Model Performance (R^2)



Agenda

- What is Machine Learning
- **Supervised Learning**
 - Regression
 - Classification
- Unsupervised Learning
 - Clustering
 - Dimension Reduction
- Model Selection & Evaluation
 - Cross Validation
 - Hyperparameter Tuning
- Intro to Deep Learning

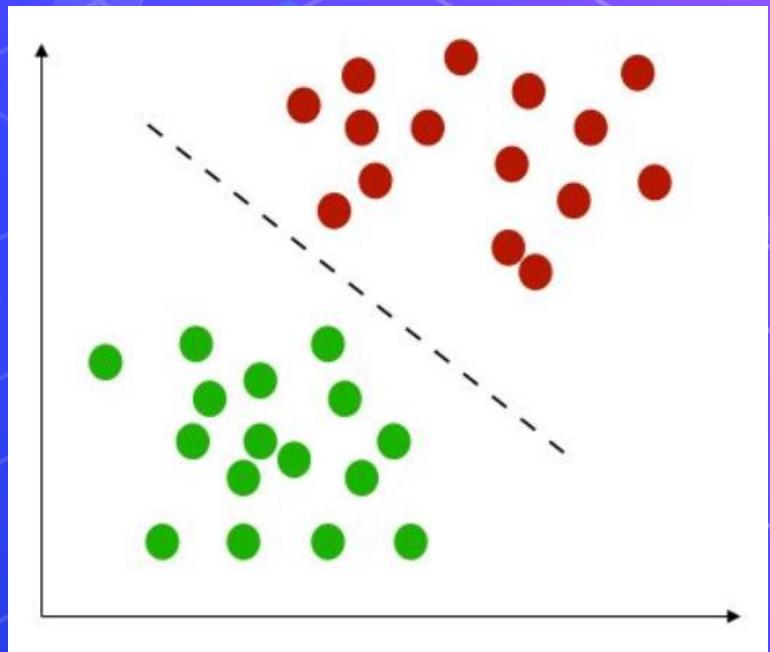


Classification

Classification specifies the class to which data elements belong to and is best used when the output has finite and discrete values. It predicts a class for an input variable as well.

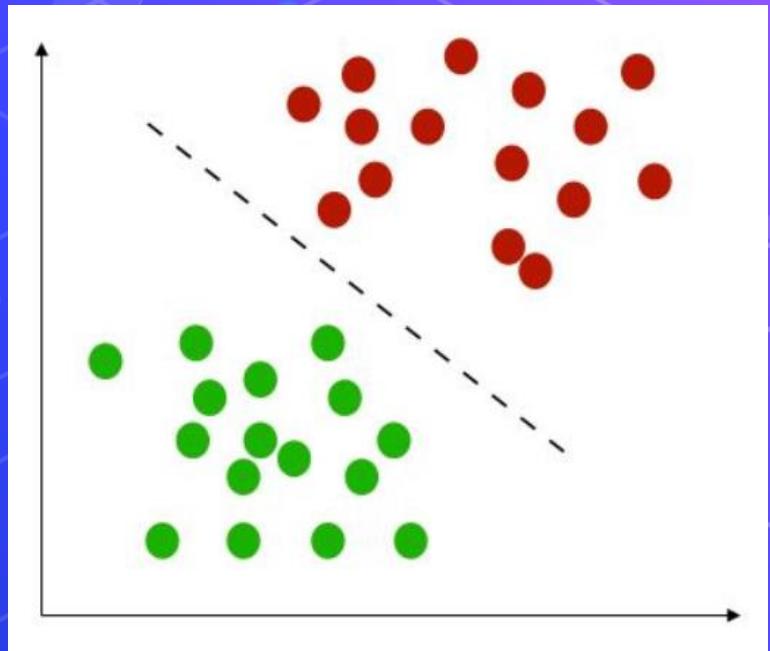
For example,
it can be used to classify the animal in an image, it's either a cat or a dog.

- Email spam detection
- Face classification
- Patient has cancer or not
- Fraud detection
- ...



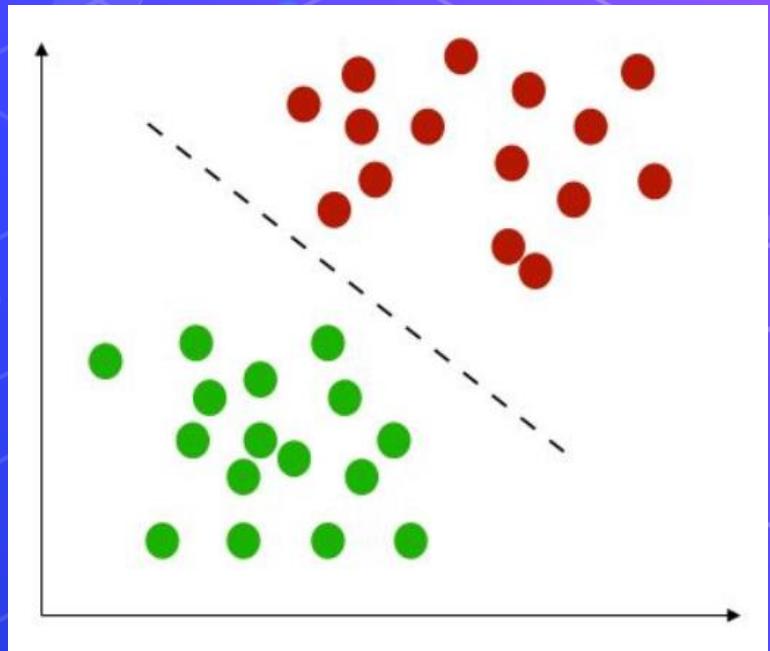
Classification

- Logistic Regression
- K-Nearest Neighbors (KNN)
- Naive Bayes
- SVM
- Decision Trees
- Ensemble Methods
 - What is Bagging & Boosting
 - Random Forests
 - XGBoost
- Evaluating Model Performance



Classification

- Logistic Regression
- K-Nearest Neighbors (KNN)
- Naive Bayes
- SVM
- Decision Trees
- Ensemble Methods
 - What is Bagging & Boosting
 - Random Forests
 - XGBoost
- Evaluating Model Performance



Logistic Regression

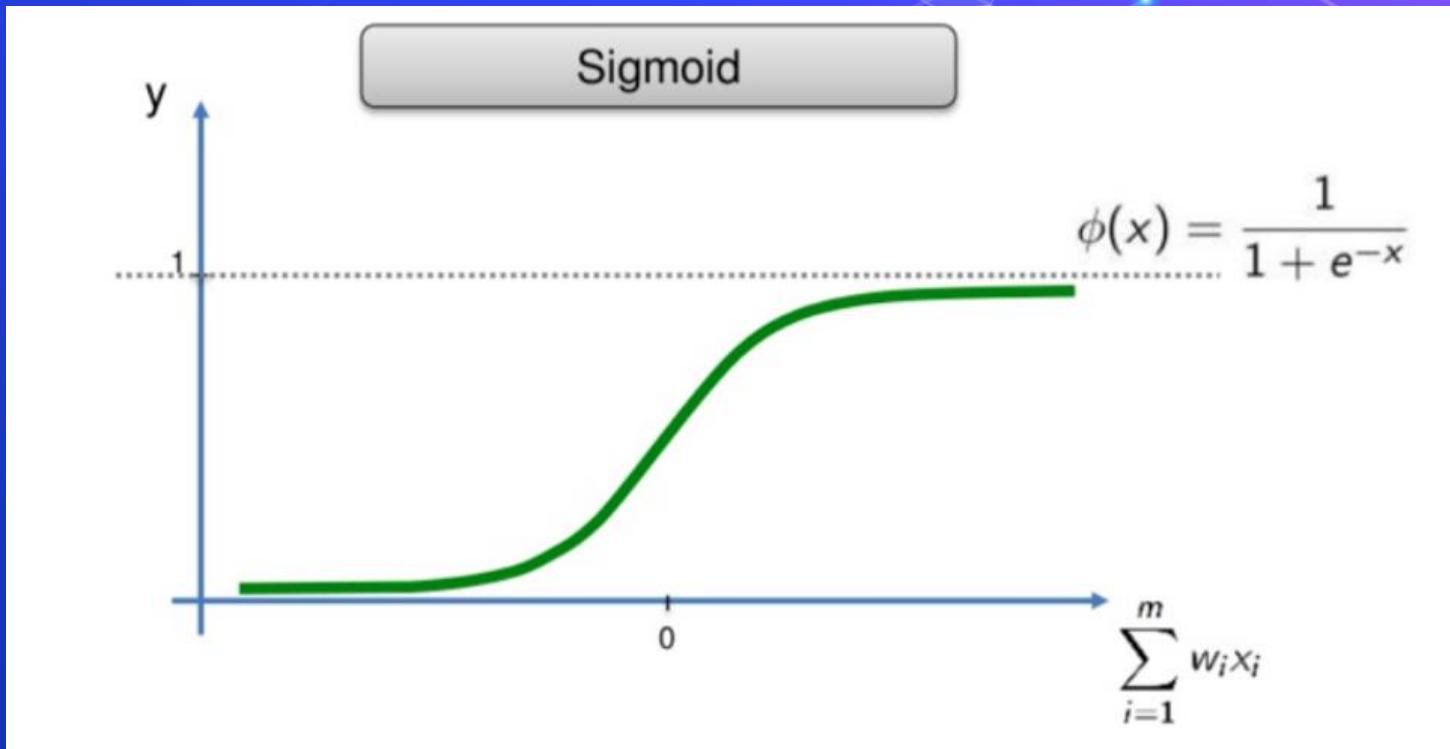
The Logistic Function

$$y = \frac{1}{1 + e^{-f(x_1, x_2, \dots, x_n)}} \in (0, 1)$$

where

$$f(x_1, x_2, \dots, x_n) = a_0 + a_1 x_1 + \dots + a_n x_n \in (-\infty, +\infty)$$

Logistic Regression



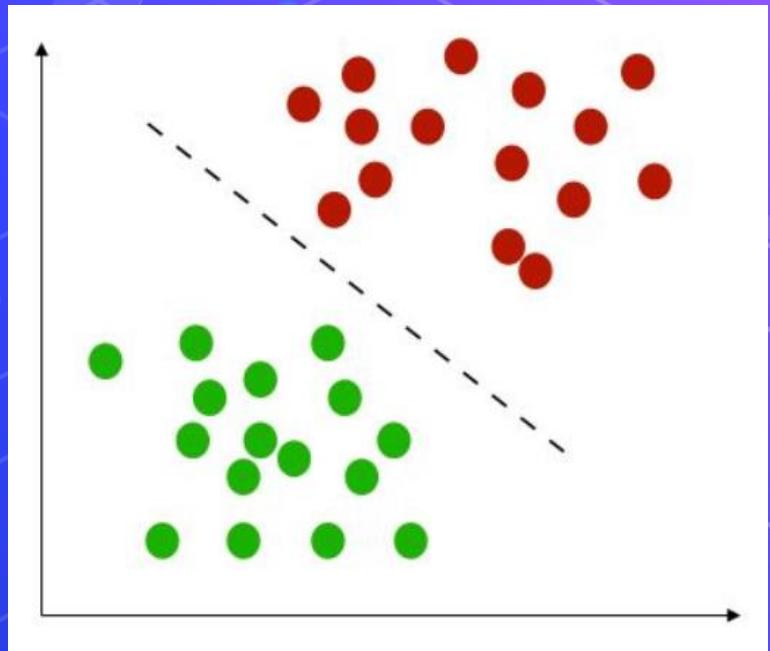
Logistic Regression

```
1 from sklearn.linear_model import LogisticRegression  
2 from sklearn.metrics import classification_report  
3 from sklearn.metrics import accuracy_score  
4 from sklearn.metrics import confusion_matrix  
5  
6  
7 model = LogisticRegression()  
8 model.fit(x_train, y_train)  
9 y_pred = model.predict(x_test)  
10  
11 print(confusion_matrix(y_test, y_pred))  
12 print(accuracy_score(y_test, y_pred))  
13 print(classification_report(y_test, y_pred))
```




Classification

- Logistic Regression
- K-Nearest Neighbors (KNN)
- Naive Bayes
- SVM
- Decision Trees
- Ensemble Methods
 - What is Bagging & Boosting
 - Random Forests
 - XGBoost
- Evaluating Model Performance



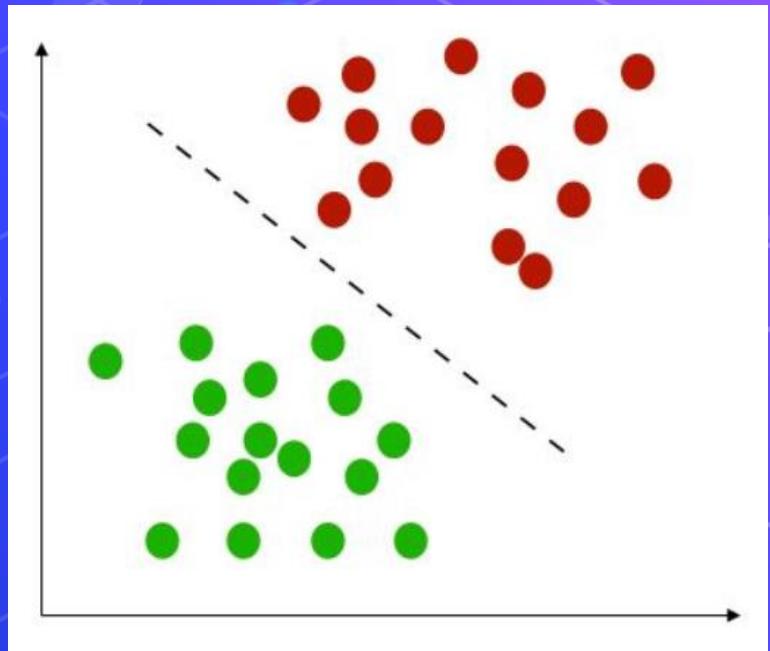
K-Nearest Neighbors (KNN)

```
● ● ●  
1 from sklearn.neighbors import KNeighborsClassifier  
2 from sklearn.metrics import classification_report  
3 from sklearn.metrics import accuracy_score  
4 from sklearn.metrics import confusion_matrix  
5  
6  
7 model = KNeighborsClassifier()  
8 model.fit(x_train, y_train)  
9 y_pred = model.predict(x_test)  
10  
11 print(confusion_matrix(y_test, y_pred))  
12 print(accuracy_score(y_test, y_pred))  
13 print(classification_report(y_test, y_pred))
```

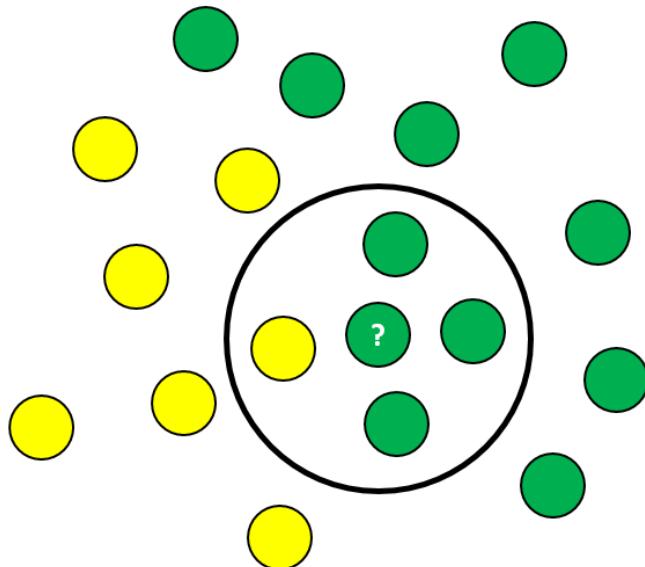


Classification

- Logistic Regression
- K-Nearest Neighbors (KNN)
- **Naive Bayes**
- SVM
- Decision Trees
- Ensemble Methods
 - What is Bagging & Boosting
 - Random Forests
 - XGBoost
- Evaluating Model Performance



Naive Bayes



$$P(\text{yellow}) = \frac{7}{17}$$

$$P(\text{green}) = \frac{10}{17}$$

$$P'(? | \text{green}) = \frac{3}{10}$$

$$P'(? | \text{yellow}) = \frac{1}{7}$$

prior probabilities

number of samples in a given class
divided by the total number of samples

we consider just the vicinity
of the new sample we want to classify

$$P''(\text{? is green}) = P(\text{green}) * P'(\text{?} | \text{green}) = \frac{10}{17} * \frac{3}{10} = \frac{30}{170}$$

$$P''(\text{? is yellow}) = P(\text{yellow}) * P'(\text{?} | \text{yellow}) = \frac{7}{17} * \frac{1}{7} = \frac{7}{119}$$

Naive Bayes

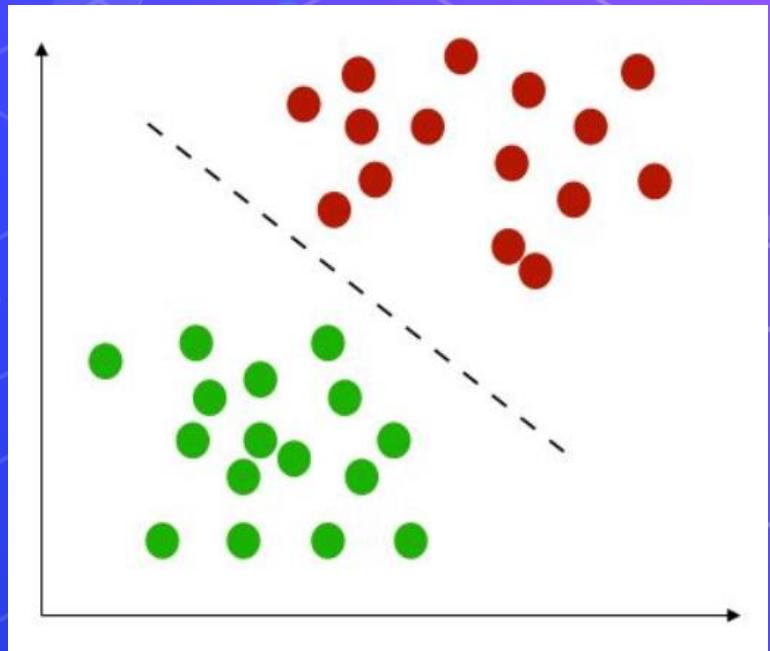


```
1 from sklearn.naive_bayes import MultinomialNB
2 from sklearn.metrics import classification_report
3 from sklearn.metrics import accuracy_score
4 from sklearn.metrics import confusion_matrix
5
6
7 model = MultinomialNB()
8 model.fit(x_train, y_train)
9 y_pred = model.predict(x_test)
10
11 print(confusion_matrix(y_test, y_pred))
12 print(accuracy_score(y_test, y_pred))
13 print(classification_report(y_test, y_pred))
```

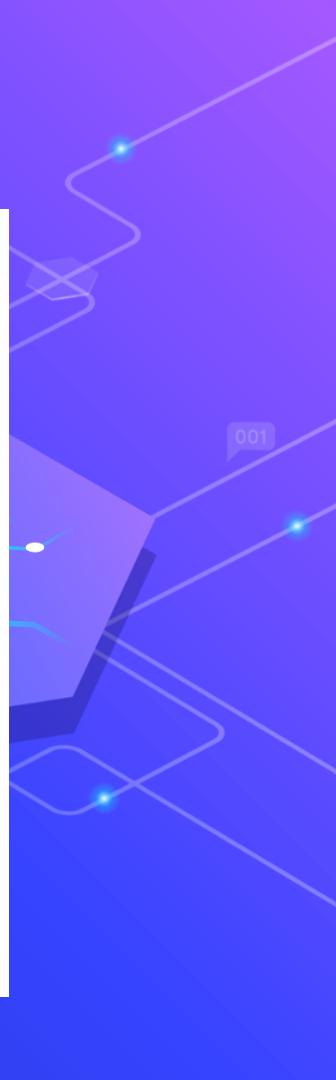
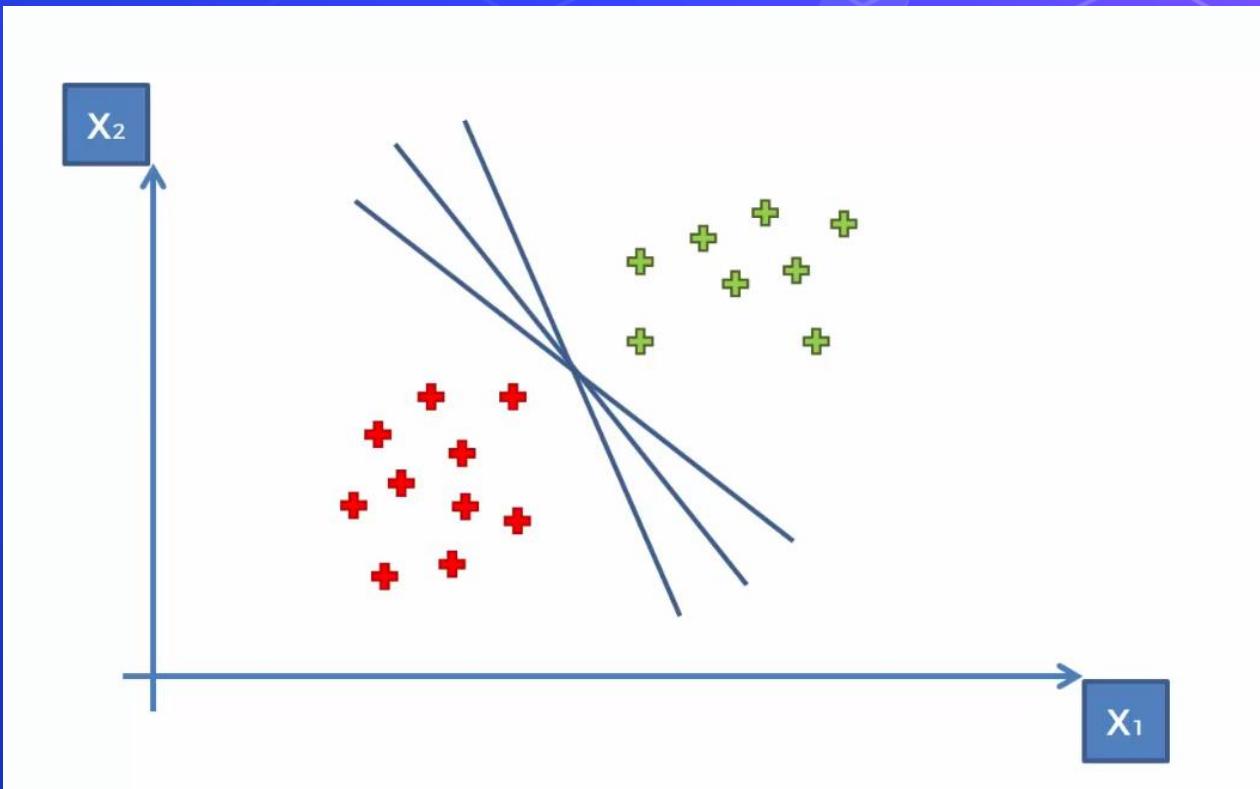


Classification


- Logistic Regression
- K-Nearest Neighbors (KNN)
- Naive Bayes
- **SVM**
- Decision Trees
- Ensemble Methods
 - What is Bagging & Boosting
 - Random Forests
 - XGBoost
- Evaluating Model Performance



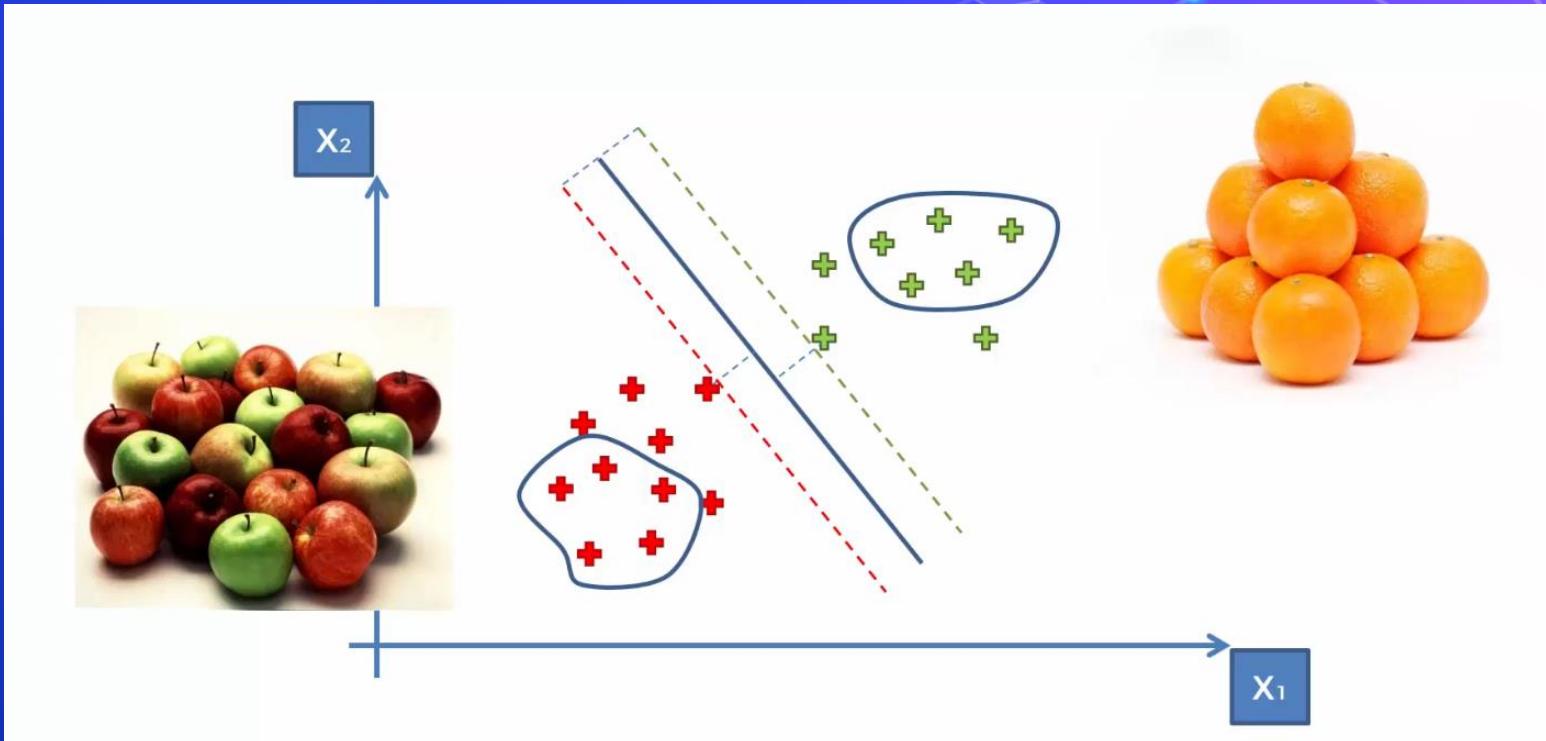
SVM (Support Vector Machine)



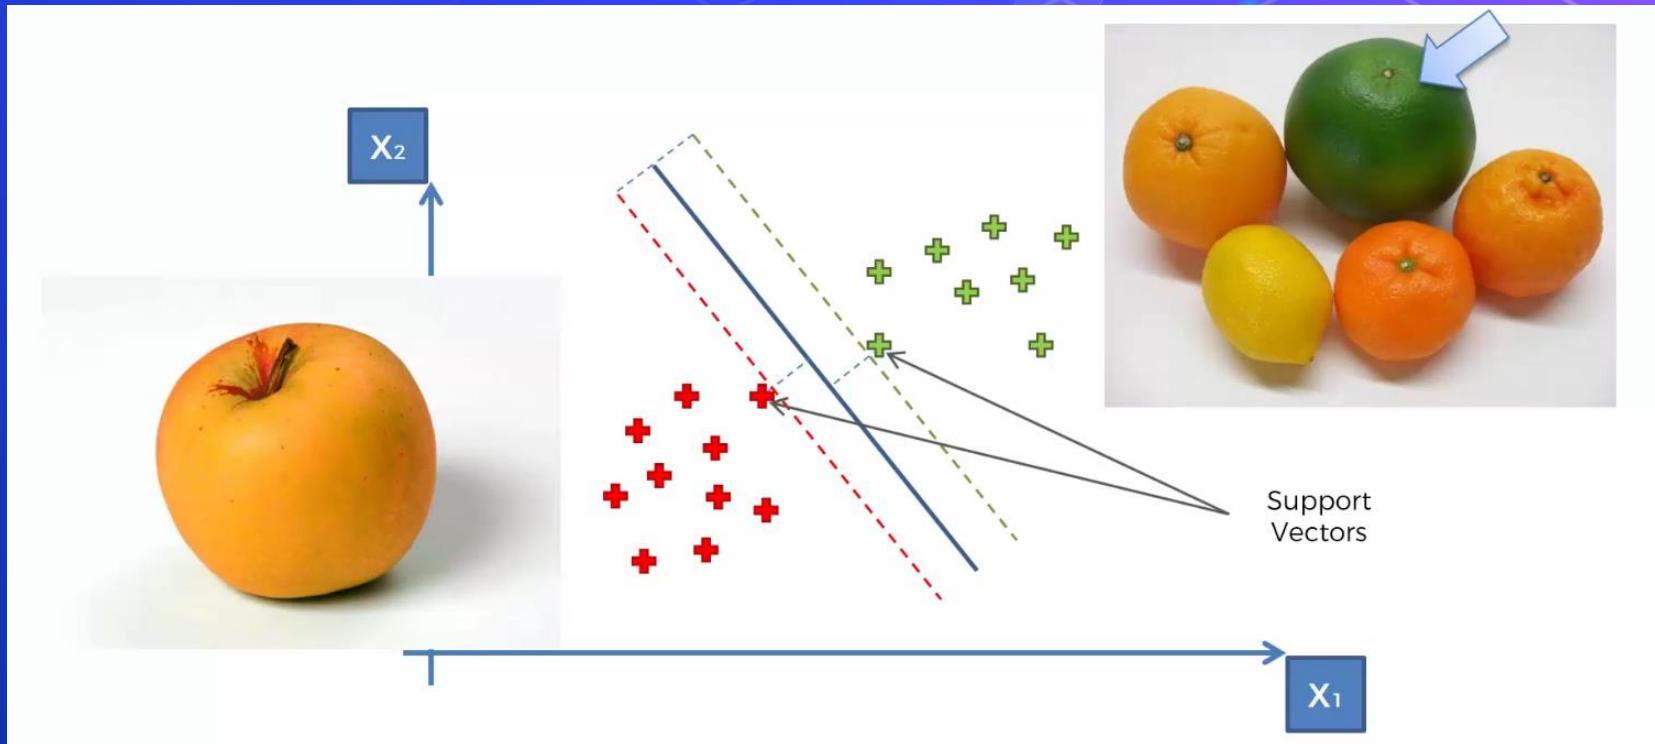
SVM (Support Vector Machine)



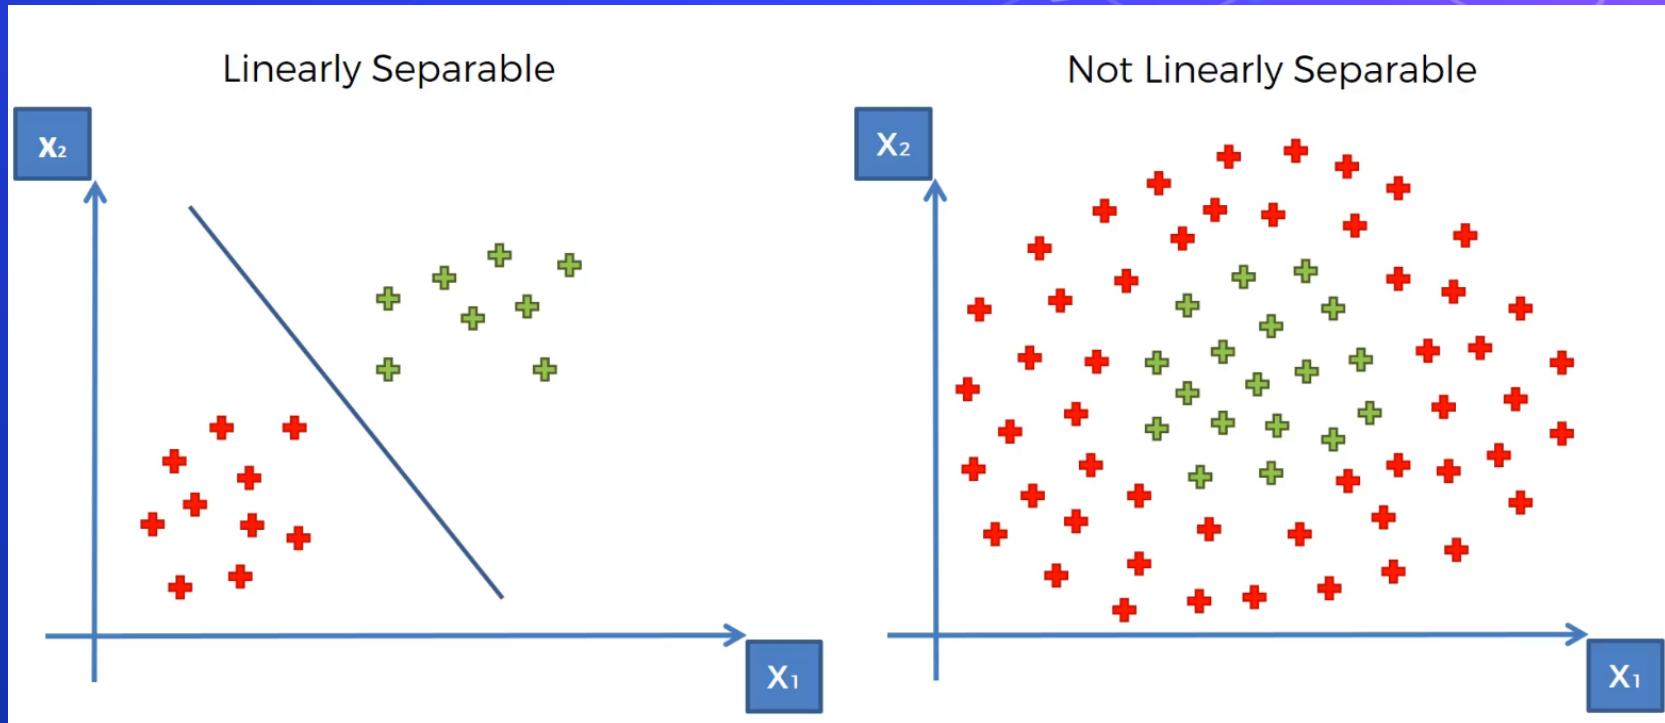
SVM (Support Vector Machine)



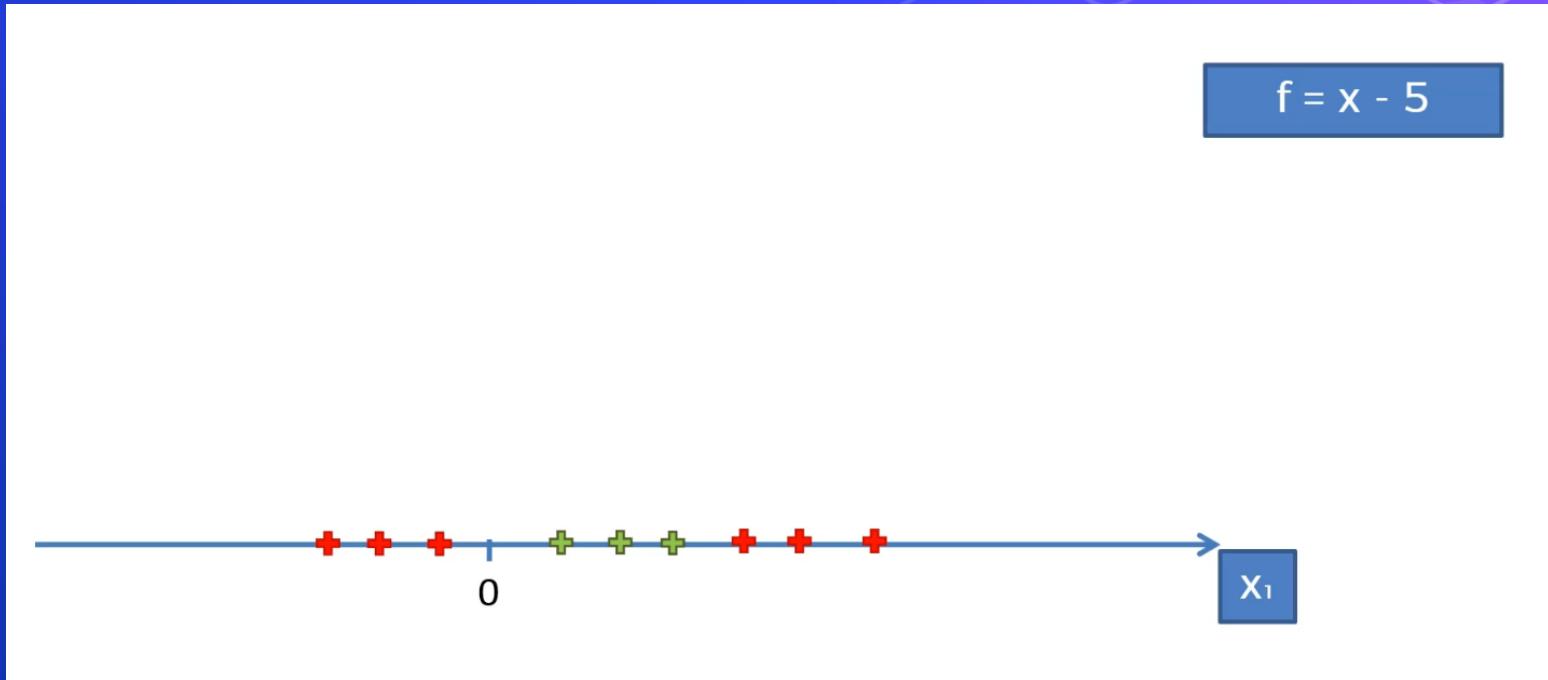
SVM (Support Vector Machine)



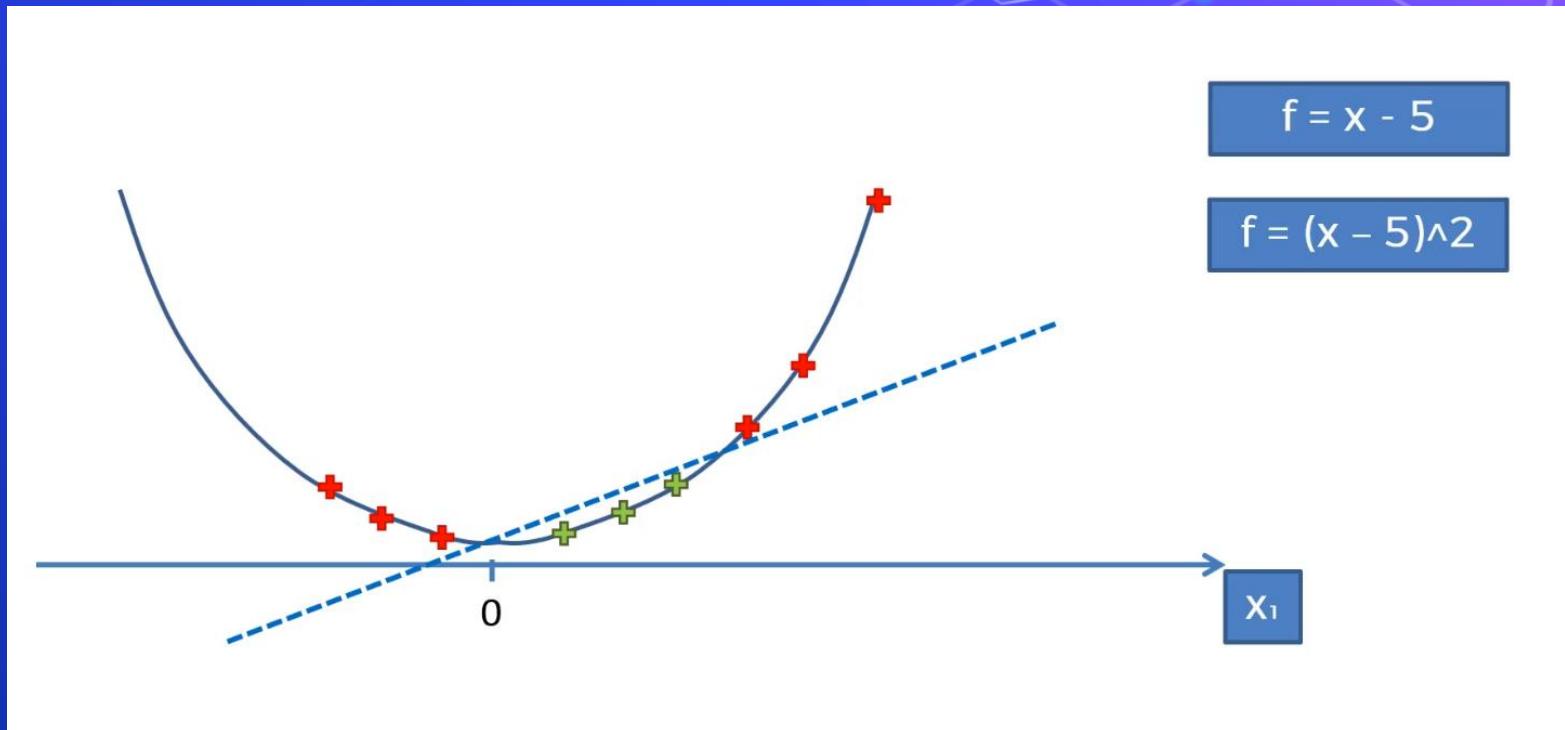
SVM (Support Vector Machine)



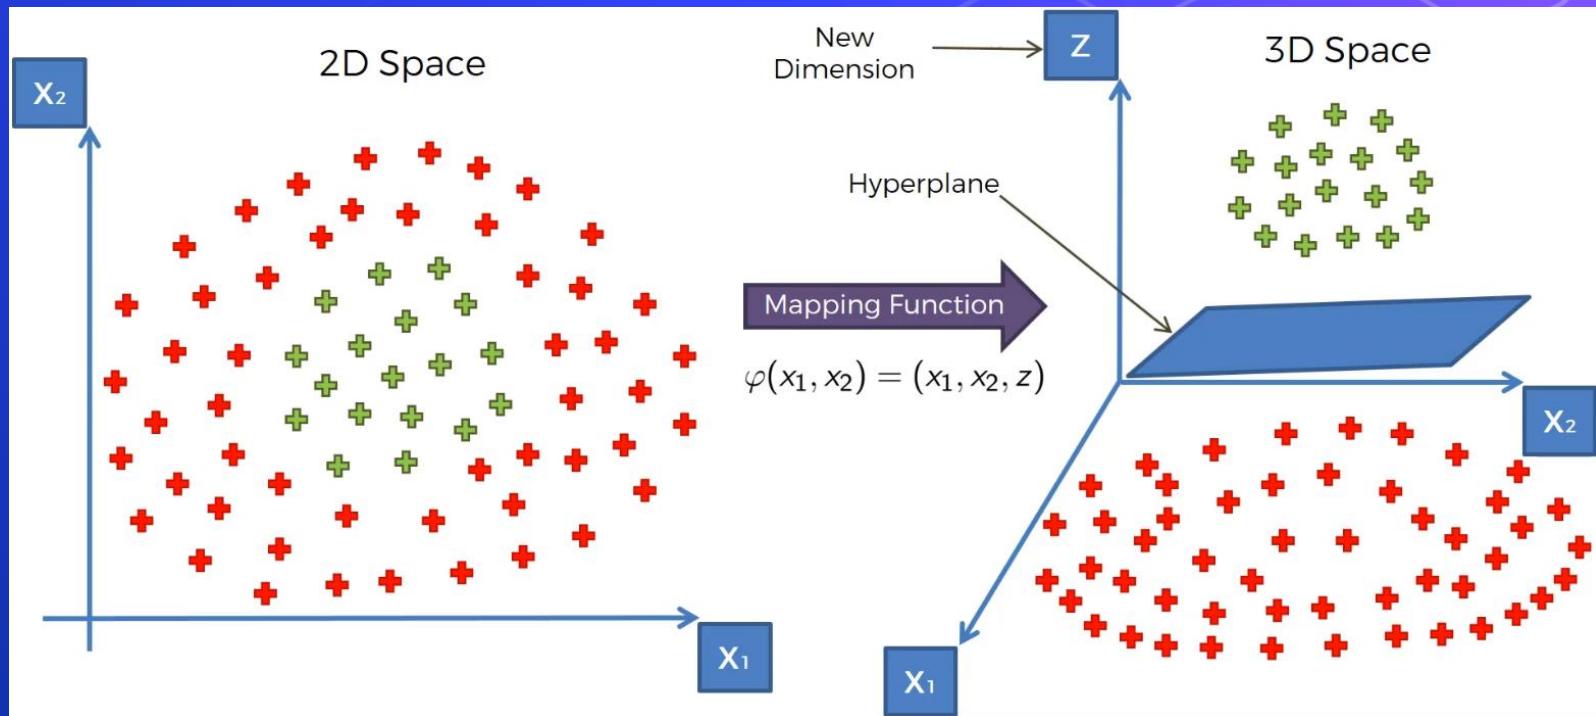
SVM (Support Vector Machine)



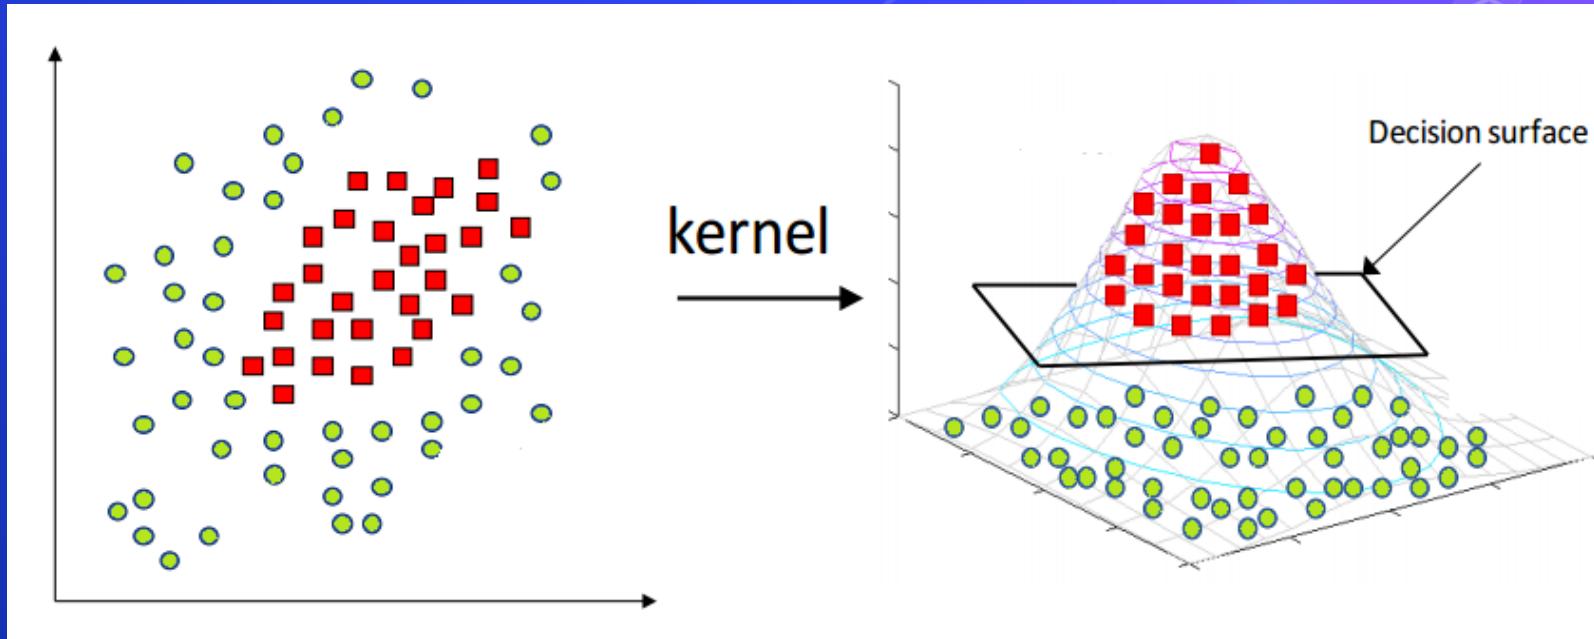
SVM (Support Vector Machine)



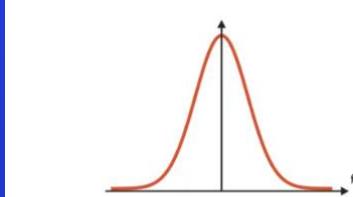
SVM (Support Vector Machine)



SVM (Support Vector Machine)




SVM (Support Vector Machine)




Gaussian RBF Kernel

$$K(\vec{x}, \vec{l}^i) = e^{-\frac{\|\vec{x}-\vec{l}^i\|^2}{2\sigma^2}}$$



Sigmoid Kernel

$$K(X, Y) = \tanh(\gamma \cdot X^T Y + r)$$



Polynomial Kernel

$$K(X, Y) = (\gamma \cdot X^T Y + r)^d, \gamma > 0$$

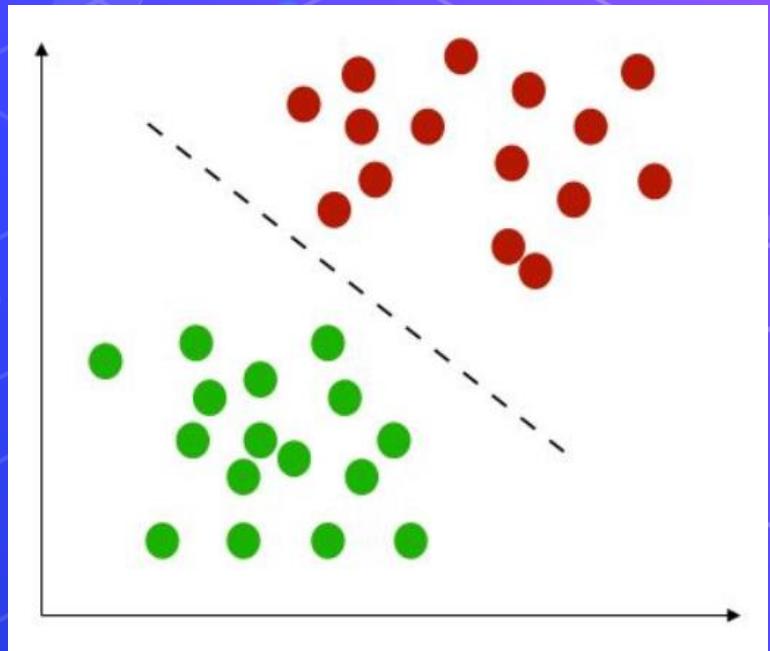
SVM (Support Vector Machine)

```
1 from sklearn.svm import SVC
2 from sklearn.metrics import classification_report
3 from sklearn.metrics import accuracy_score
4 from sklearn.metrics import confusion_matrix
5
6
7 model = SVC()
8 model.fit(x_train, y_train)
9 y_pred = model.predict(x_test)
10
11 print(confusion_matrix(y_test, y_pred))
12 print(accuracy_score(y_test, y_pred))
13 print(classification_report(y_test, y_pred))
```

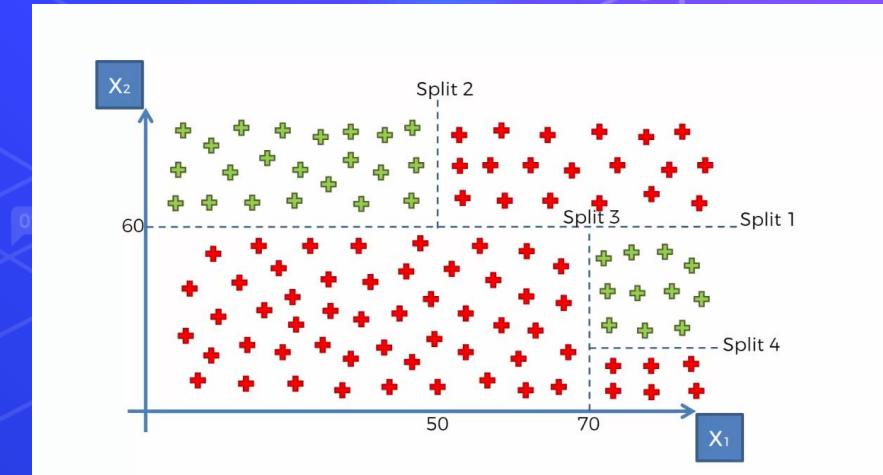



Classification

- Logistic Regression
- K-Nearest Neighbors (KNN)
- Naive Bayes
- SVM
- **Decision Trees**
- Ensemble Methods
 - What is Bagging & Boosting
 - Random Forests
 - XGBoost
- Evaluating Model Performance



Decision Trees



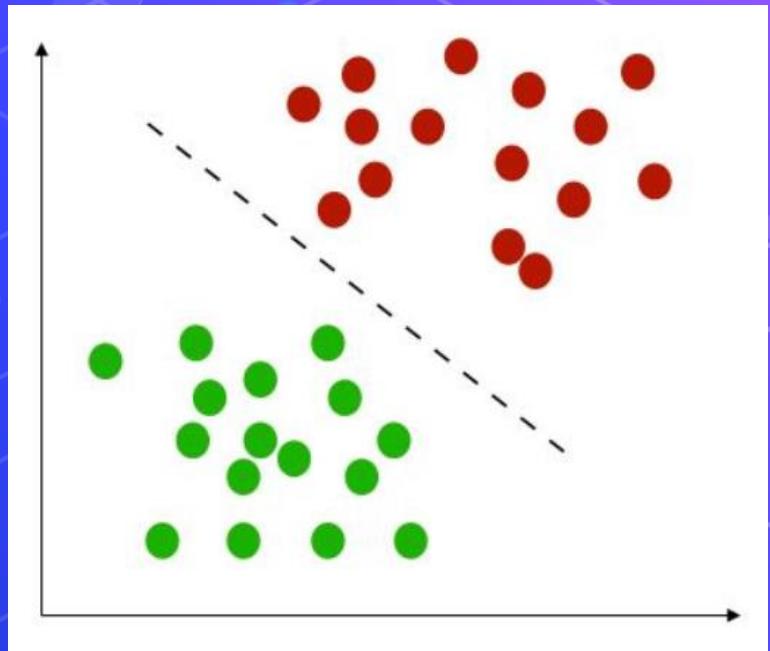
Decision Trees

```
1 from sklearn.tree import DecisionTreeClassifier  
2 from sklearn.metrics import classification_report  
3 from sklearn.metrics import accuracy_score  
4 from sklearn.metrics import confusion_matrix  
5  
6  
7 model = DecisionTreeClassifier()  
8 model.fit(x_train, y_train)  
9 y_pred = model.predict(x_test)  
10  
11 print(confusion_matrix(y_test, y_pred))  
12 print(accuracy_score(y_test, y_pred))  
13 print(classification_report(y_test, y_pred))
```

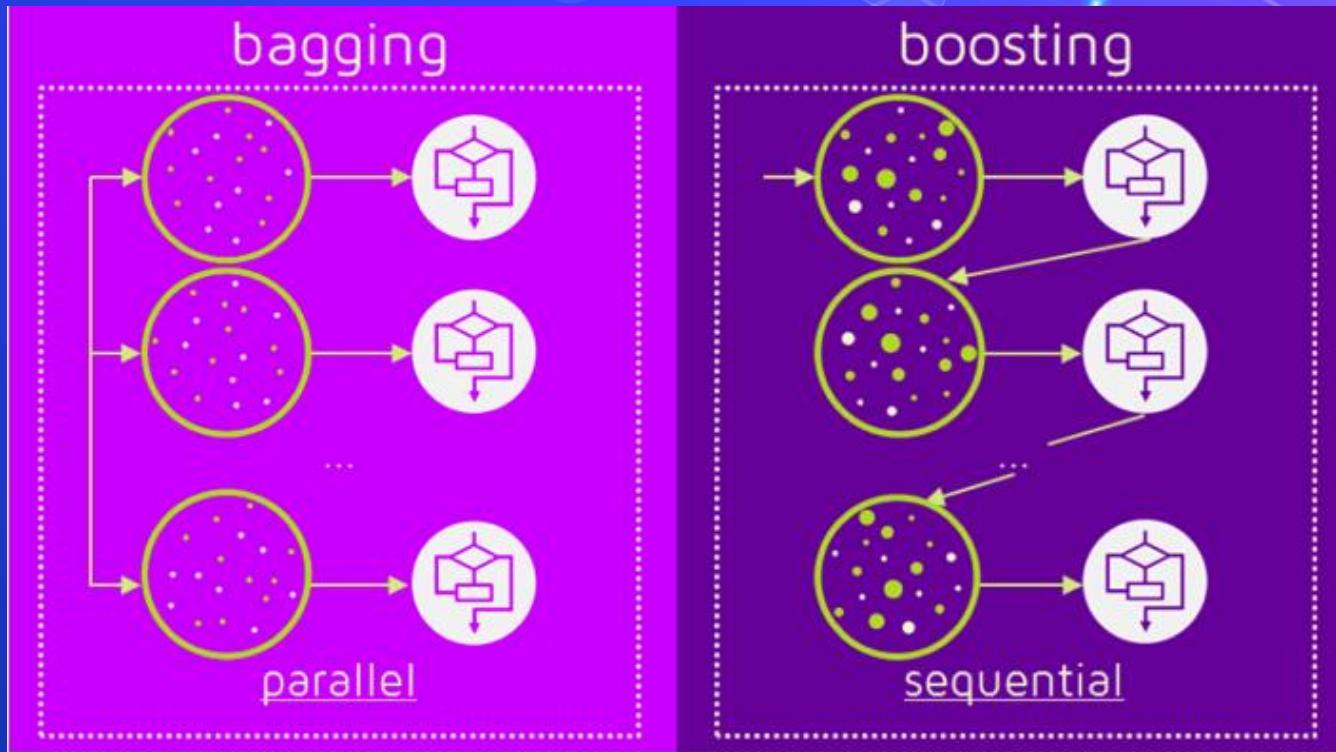


Classification

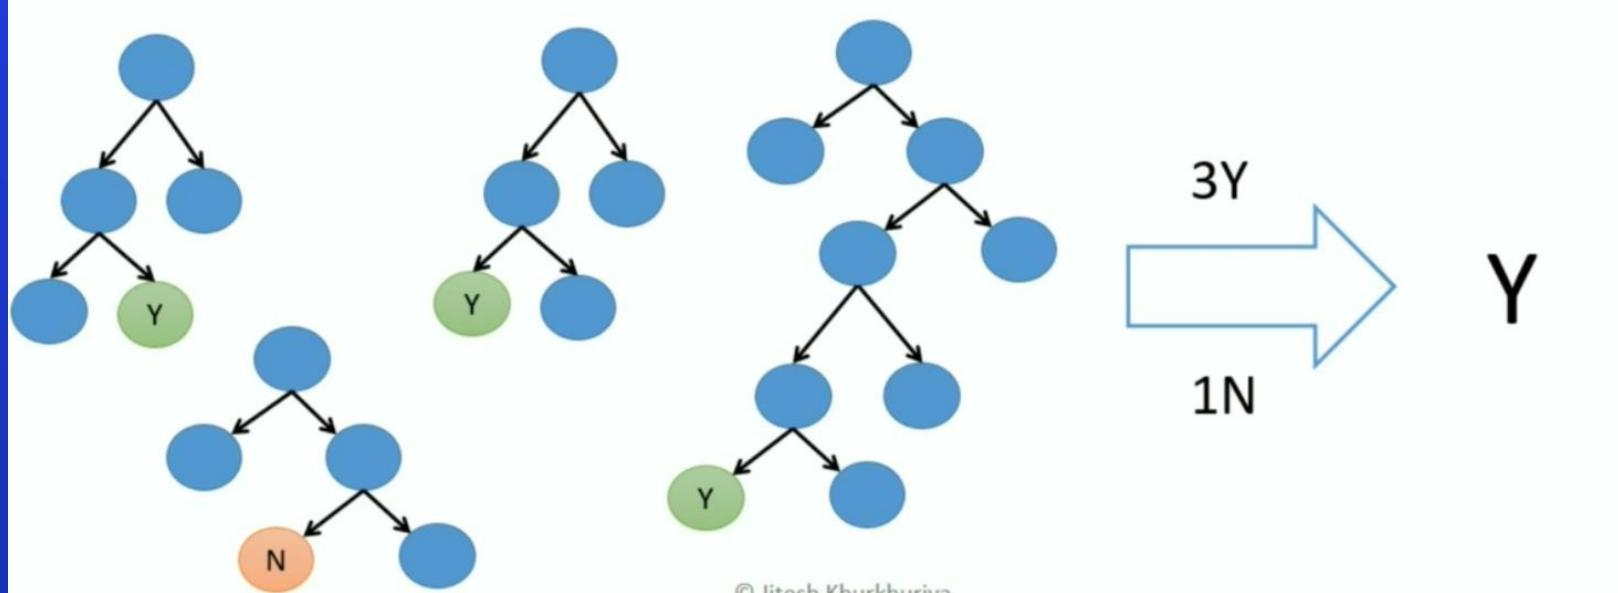
- Logistic Regression
- K-Nearest Neighbors (KNN)
- Naive Bayes
- SVM
- Decision Trees
- **Ensemble Methods**
 - What is Bagging & Boosting
 - Random Forests
 - XGBoost
- Evaluating Model Performance



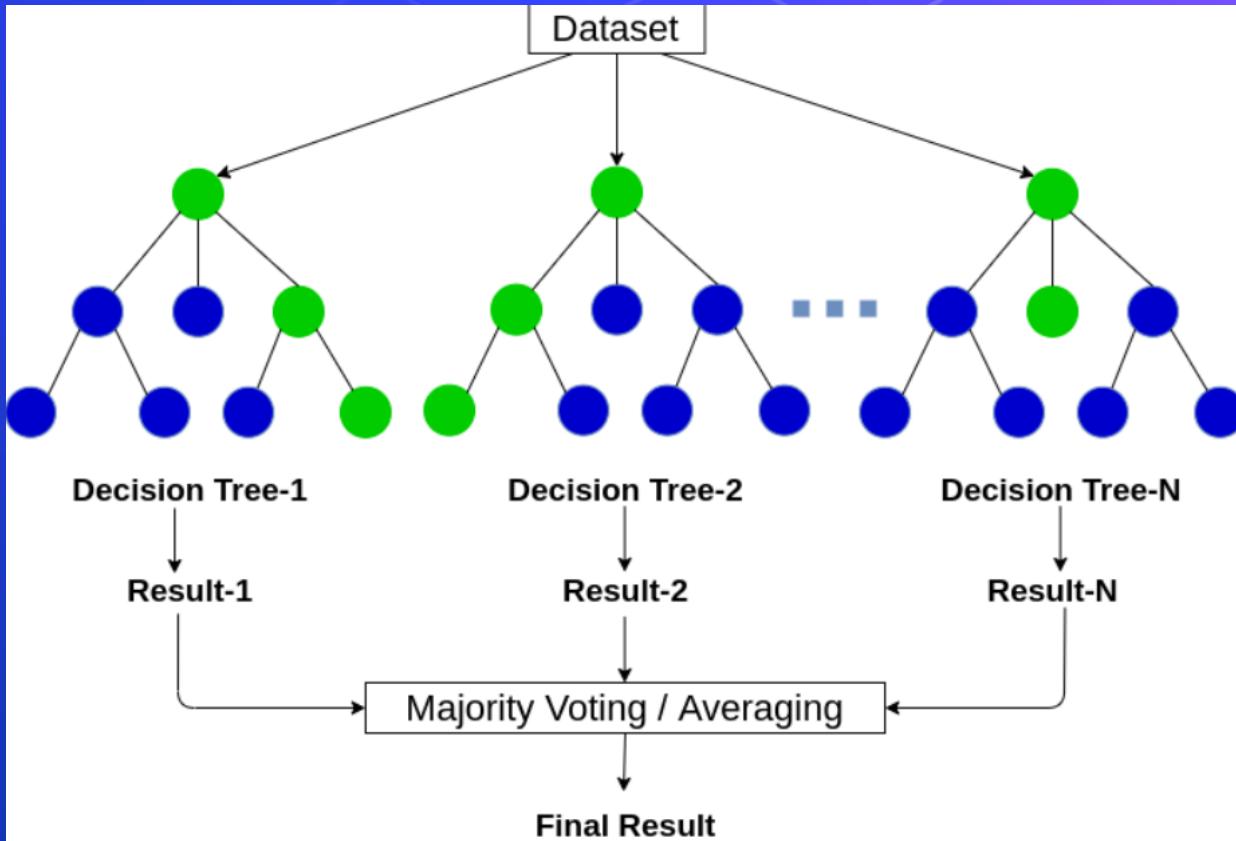
What is Bagging & Boosting



What is Bagging & Boosting



Random Forests




Random Forests

```
1 from sklearn.ensemble import RandomForestClassifier  
2 from sklearn.metrics import classification_report  
3 from sklearn.metrics import accuracy_score  
4 from sklearn.metrics import confusion_matrix  
5  
6  
7 model = RandomForestClassifier()  
8 model.fit(x_train, y_train)  
9 y_pred = model.predict(x_test)  
10  
11 print(confusion_matrix(y_test, y_pred))  
12 print(accuracy_score(y_test, y_pred))  
13 print(classification_report(y_test, y_pred))
```



What is Bagging & Boosting



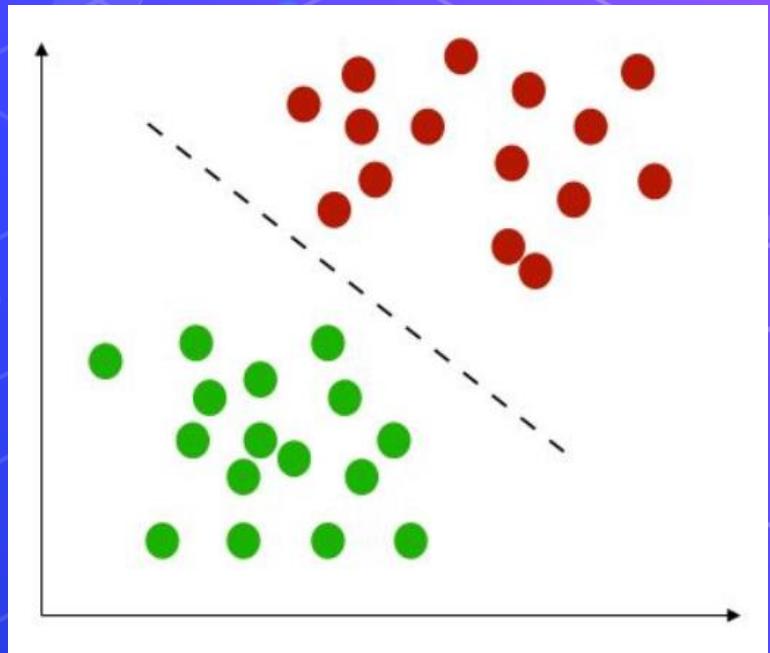
XGBoost

```
1 import xgboost as xgb
2 from sklearn.metrics import classification_report
3 from sklearn.metrics import accuracy_score
4 from sklearn.metrics import confusion_matrix
5
6
7 model = xgb.XGBClassifier()
8 model.fit(x_train, y_train)
9 y_pred = model.predict(x_test)
10
11 print(confusion_matrix(y_test, y_pred))
12 print(accuracy_score(y_test, y_pred))
13 print(classification_report(y_test, y_pred))
```



Classification

- Logistic Regression
- K-Nearest Neighbors (KNN)
- Naive Bayes
- SVM
- Decision Trees
- Ensemble Methods
 - What is Bagging & Boosting
 - Random Forests
 - XGBoost
- Evaluating Model Performance




Evaluating Model Performance (Confusion Matrix)

		Actual Value	
		Positive	Negative
Predicted Value	Positive	TP (True Positive)	FP (False Positive)
	Negative	FN (False Negative)	TN (True Negative)

- True Positive (TP) : Observation is positive, and is predicted to be positive.
- False Negative (FN) : Observation is positive, but is predicted negative.
- True Negative (TN) : Observation is negative, and is predicted to be negative.
- False Positive (FP) : Observation is negative, but is predicted positive.

Evaluating Model Performance (Confusion Matrix)



True positives (TP): actuals are positives and are predicted as positives.

You predicted that a woman is pregnant and she actually is.

False positives (FP) - Type 1 Error: actuals are negatives and are predicted as positives.

You predicted that a man is pregnant but he actually is not.

False negatives (FN) - Type 2 Error: actuals are positives and are predicted as negatives.

You predicted that a woman is not pregnant but she actually is.

True negatives (TN): actuals are negatives and are predicted as positives.

You predicted that a man is not pregnant and he actually is not.

Evaluating Model Performance (Confusion Matrix)

```
1 from sklearn.metrics import confusion_matrix  
2  
3  
4 confusion_matrix(y_pred, y_true)  
5  
6 """  
7 array([[43,  2],  
8        [11, 87]], dtype=int64)  
9 """
```



Evaluating Model Performance (Accuracy)

Accuracy

- It's a measure of how good the model is.
- It's the ratio of the true predicted values over all the values.

Accuracy:

$$ACC = \frac{TP + TN}{TP + TN + FP + FN}$$

Recall:

$$Recall = \frac{TP}{TP + FN}$$

Precision:

$$Precision = \frac{TP}{TP + FP}$$

F₁ score:

$$F_1 = \frac{2}{\frac{1}{Recall} + \frac{1}{Precision}}$$

Evaluating Model Performance (Accuracy)

```
1 from sklearn.metrics import accuracy_score  
2  
3  
4 accuracy_score(y_pred, y_true)  
5  
6 """  
7 0.9090909090909091  
8 """
```



Evaluating Model Performance (Recall or Sensitivity)

- TP will be that the COVID 19 residents diagnosed with COVID-19.
- TN will be that healthy residents are with good health.
- FP will be that those actually healthy residents are predicted as COVID 19 residents.
- FN will be that those actual COVID 19 residents are predicted as the healthy residents.

which scenario do you think will have the highest cost?

Imagine that if we predict COVID-19 residents as healthy patients and they do not need to quarantine, there would be a massive number of COVID-19 infections. The cost of false negatives is much higher than the cost of false positives.

Case 1

COVID 19 = 1
Healthy = 0

Cost of FN > Cost of FP

Actual

Healthy predicted as sick

Sick predicted as healthy

Predict	COVID 19 (1)	Diagnosed COVID 19 (1)	Diagnosed Healthy (0)
Covid 19 (1)	TP	FP	
Healthy (0)	FN	TN	

Evaluating Model Performance (Recall or Sensitivity)

Recall (Sensitivity)

- It's a ratio of True Positives to all the positives in your data.
- Low recall: the more False Negatives the model predicts, the lower the recall.

$$\text{Recall} = \frac{\text{True Positive}}{\text{Predicted Results}} \quad \text{or} \quad \frac{\text{True Positive}}{\text{True Positive} + \text{False Negative}}$$

Accuracy:

$$ACC = \frac{TP + TN}{TP + TN + FP + FN}$$

Recall:

$$\text{Recall} = \frac{TP}{TP + FN}$$

Precision:

$$\text{Precision} = \frac{TP}{TP + FP}$$

F₁ score:

$$F_1 = \frac{2}{\frac{1}{\text{Recall}} + \frac{1}{\text{Precision}}}$$

Evaluating Model Performance (Recall or Sensitivity)

```
1 from sklearn.metrics import recall_score  
2  
3  
4 recall_score(y_pred, y_true)  
5  
6 """  
7 0.9090909090909091  
8 """
```



Evaluating Model Performance (Precision)

- TP will be that spam emails are placed in the spam folder.
- TN will be that important emails are received.
- FP will be that important emails are placed in the spam folder.
- FN will be that spam emails are received.

which scenario do you think will have the highest cost?

Well, since missing important emails will clearly be more of a problem than receiving spam, we can say that in this case, FP will have a higher cost than FN.

Case 2

Spam = 1
Not Spam = 0

Cost of FP > Cost of FN

Actual

Not spam predicted as spam

Predict

	Spam (1)	Not Spam (0)
Spam (1)	TP	FP
Not Spam (0)	FN	TN

Spam predicted as not spam

Evaluating Model Performance (Precision)

Precision

- It's a ratio of True Positives to all the positives predicted by the model.
- Low precision: the more False positives the model predicts, the lower the precision.

$$\text{Precision} = \frac{\text{True Positive}}{\text{Actual Results}} \quad \text{or} \quad \frac{\text{True Positive}}{\text{True Positive} + \text{False Positive}}$$

Accuracy:

$$ACC = \frac{TP + TN}{TP + TN + FP + FN}$$

Recall:

$$\text{Recall} = \frac{TP}{TP + FN}$$

Precision:

$$\text{Precision} = \frac{TP}{TP + FP}$$

F₁ score:

$$F_1 = \frac{2}{\frac{1}{\text{Recall}} + \frac{1}{\text{Precision}}}$$

Evaluating Model Performance (Precision)

```
1 from sklearn.metrics import precision_score  
2  
3  
4 precision_score(y_pred, y_true)  
5  
6 """  
7 0.9090909090909091  
8 """
```



Evaluating Model Performance

Case 1



COVID 19/ Healthy

Case 2



Spam/Not Spam

Cost of **FN** > Cost of **FP**

Cost of **FP** > Cost of **FN**

Recall

$$\text{recall} = \frac{\text{true positives}}{\text{true positives} + \text{false negatives}}$$

Precision

$$\text{precision} = \frac{\text{true positives}}{\text{true positives} + \text{false positives}}$$

Evaluating Model Performance (F-1 Score)

F-1 Score

- F-Score called the Harmonic mean of precision and recall.
- F-Score provides a single score that balances both the concerns of precision and recall in one number.
- A good F1 score means that you have low false positives and low false negatives.
- An F1 score is considered perfect when it's 1, while the model is a total failure when it's 0.

$$F_1 = 2 * \frac{precision * recall}{precision + recall}$$

Accuracy:

$$ACC = \frac{TP + TN}{TP + TN + FP + FN}$$

Recall:

$$Recall = \frac{TP}{TP + FN}$$

Precision:

$$Precision = \frac{TP}{TP + FP}$$

F₁ score:

$$F_1 = \frac{2}{\frac{1}{Recall} + \frac{1}{Precision}}$$

Evaluating Model Performance (F-1 Score)

```
1 from sklearn.metrics import f1_score  
2  
3  
4 f1_score(y_pred, y_true)  
5  
6 """  
7 0.9090909090909091  
8 """
```



Evaluating Model Performance (Classification Report)

```
1 from sklearn.metrics import classification_report  
2  
3  
4 classification_report(y_true, y_pred)  
5  
6 """  
7         precision    recall    f1-score    support  
8     class 0       0.50      1.00      0.67        1  
9     class 1       0.00      0.00      0.00        1  
10    class 2       1.00      0.67      0.80        3  
11 """
```

Evaluating Model Performance (F-beta Score)

F-beta Score

- Generalization of F-Score where beta is a parameter to weight the result of the precision or recall.
- Beta = 0.5 to focus more on precision, It's called F-0.5 Score.
- Beta = 2 to focus more on recall, It's called F-2 Score.
- Beta value should be used depending on your domain knowledge business case.

$$F_{\beta} = (1 + \beta^2) \cdot \frac{\text{precision} \cdot \text{recall}}{(\beta^2 \cdot \text{precision}) + \text{recall}}.$$



Evaluating Model Performance (F-beta Score)

```
1 from sklearn.metrics import fbeta_score
2
3
4 fbeta_score(y_true, y_pred, beta=0.5)
5
6 """
7 0.9090909090909091
8 """
```



Agenda

- What is Machine Learning
- Supervised Learning
 - Regression
 - Classification
- Unsupervised Learning
 - Clustering
 - Dimension Reduction
- Model Selection & Evaluation
 - Cross Validation
 - Hyperparameter Tuning
- Intro to Deep Learning




Clustering

Clustering is the task of dividing the population or data points into a number of groups such that data points in the same groups are more similar to other data points in the same group and dissimilar to the data points in other groups.

For example,
it can be used to cluster the university students to many segments depending on features like scores, departments, etc.


- University students clustering
- Mall customers segmentation
- Recommendation systems
- ...



Clustering

○ KMeans


○ Hierarchical Clustering



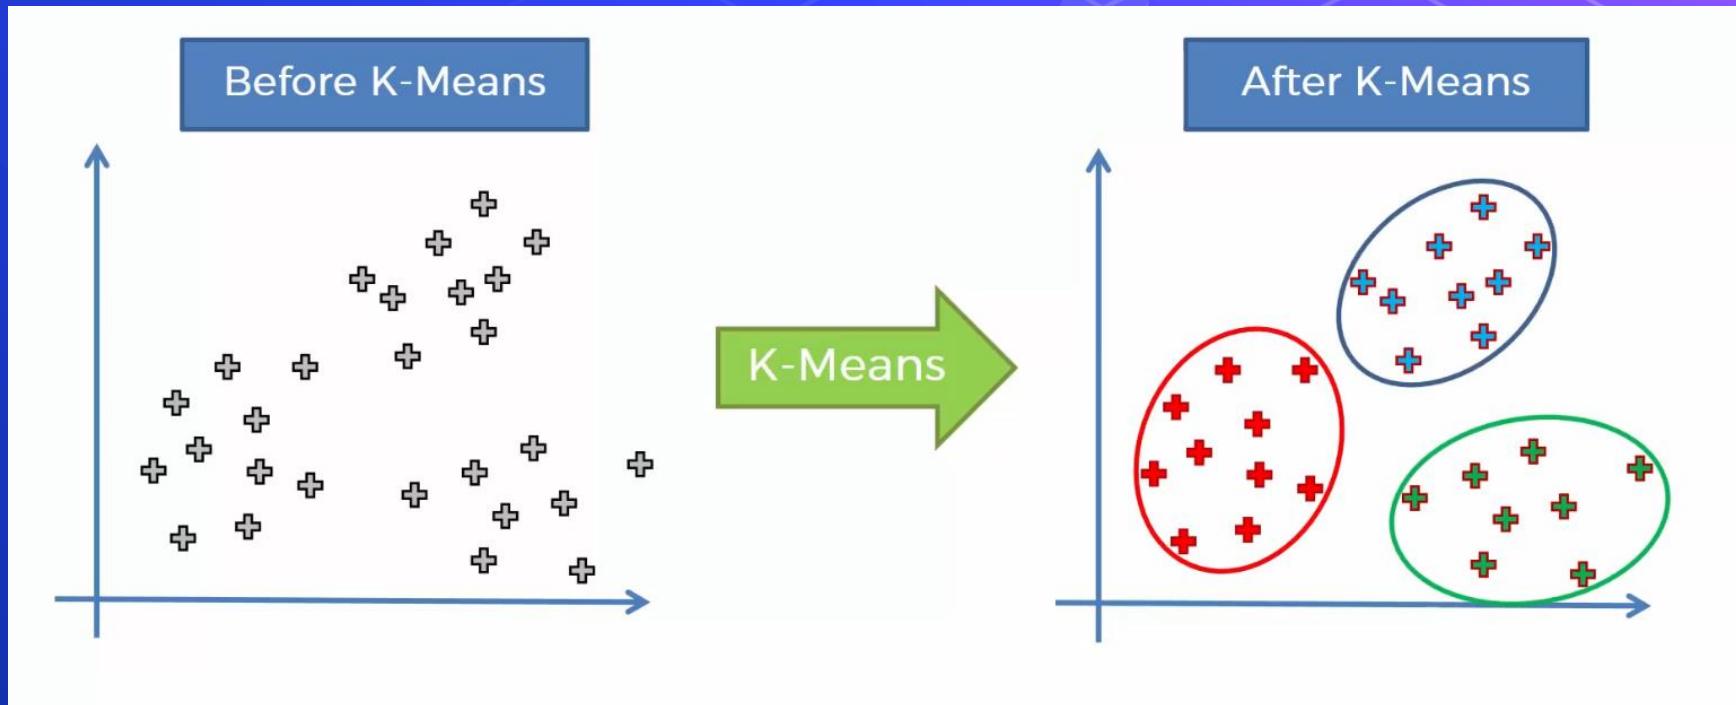
Clustering

○ KMeans

○ Hierarchical Clustering




KMeans



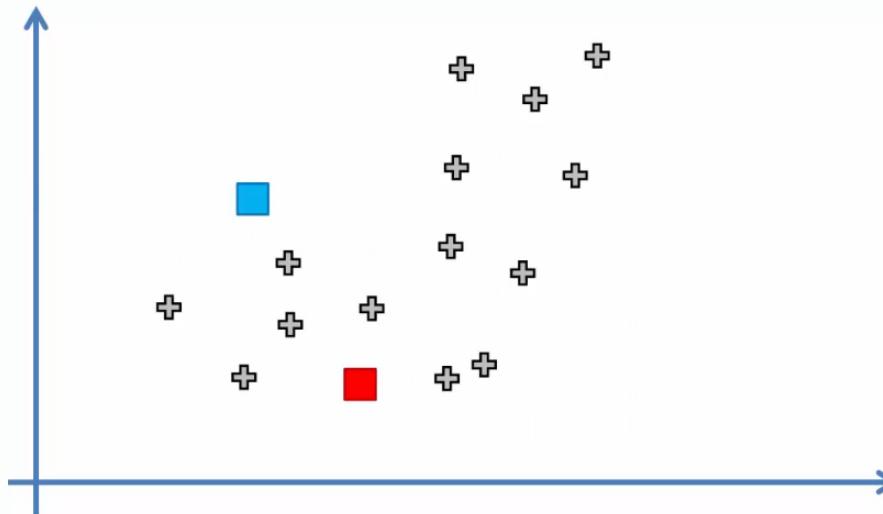
KMeans

STEP 1: Choose the number K of clusters: K = 2



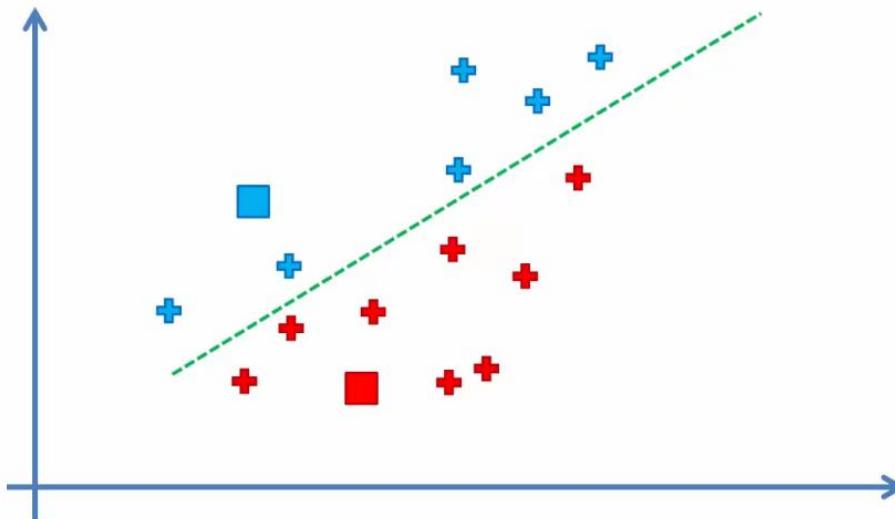
KMeans

STEP 2: Select at random K points, the centroids (not necessarily from your dataset)



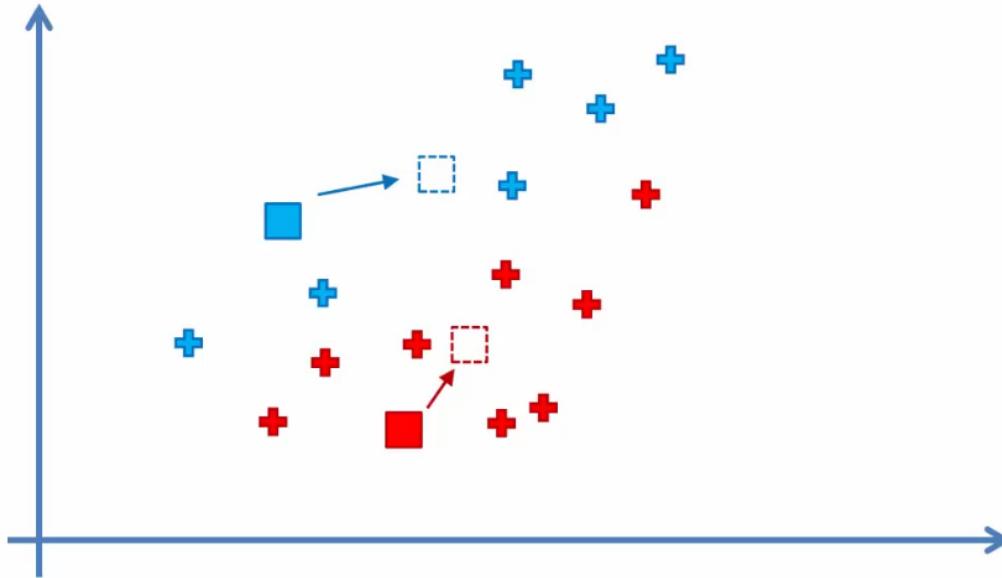
KMeans

STEP 3: Assign each data point to the closest centroid → That forms K clusters



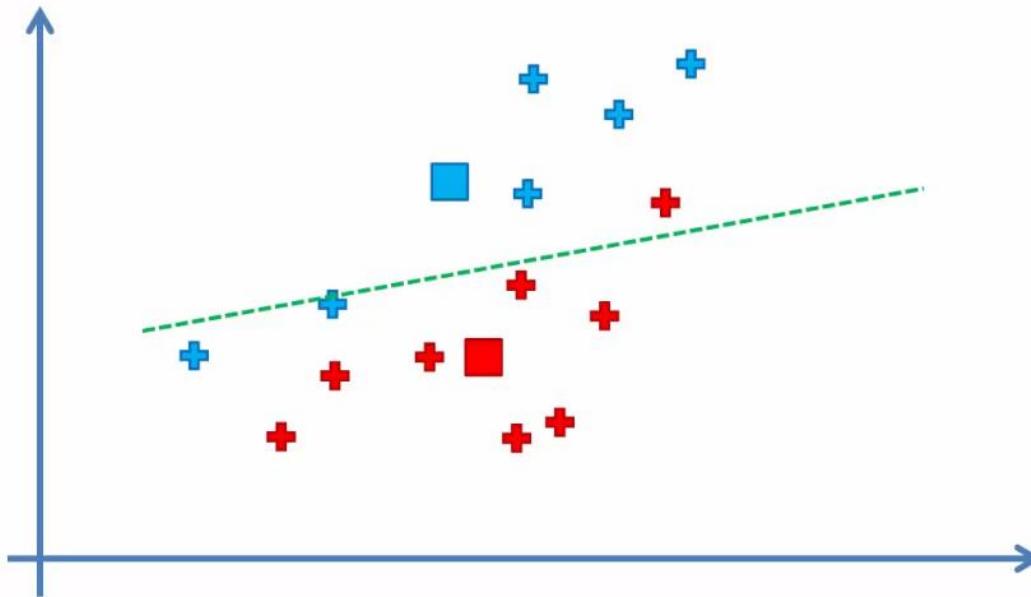
KMeans

STEP 4: Compute and place the new centroid of each cluster



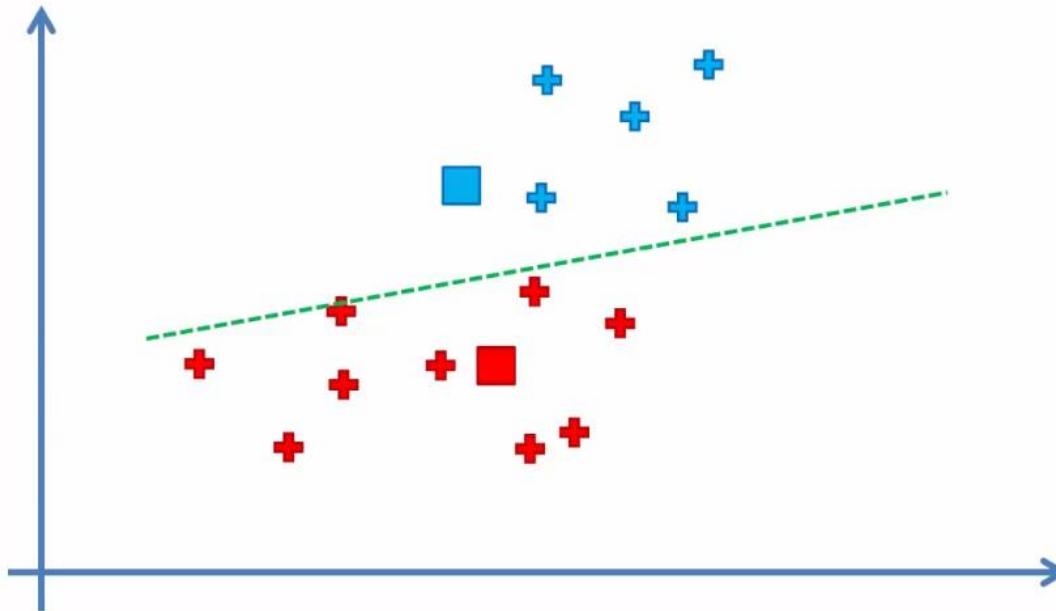
KMeans

STEP 5: Reassign each data point to the new closest centroid.
If any reassignment took place, go to STEP 4, otherwise go to FIN.



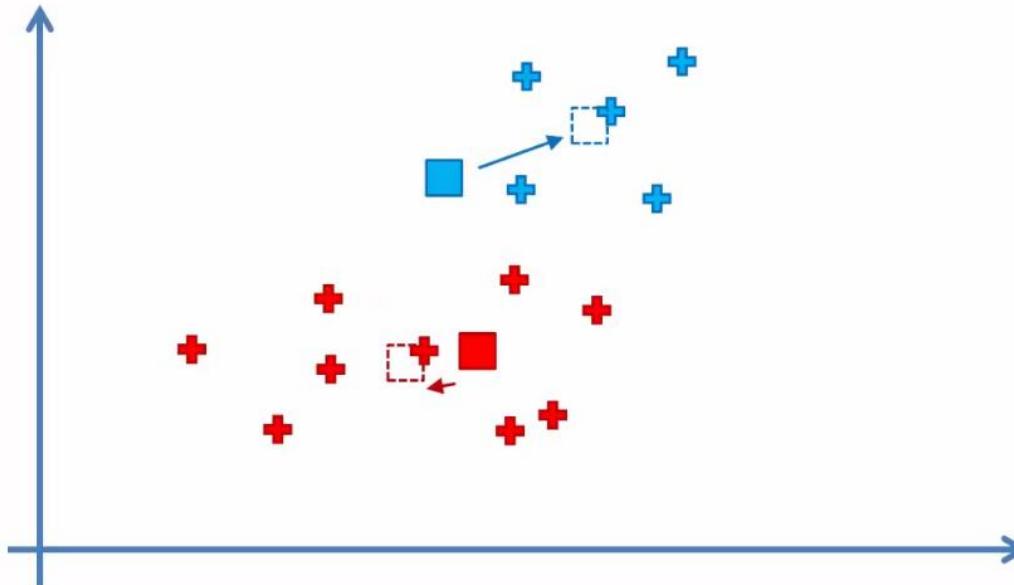
KMeans

STEP 5: Reassign each data point to the new closest centroid.
If any reassignment took place, go to STEP 4, otherwise go to FIN.




KMeans

STEP 4: Compute and place the new centroid of each cluster



KMeans

STEP 5: Reassign each data point to the new closest centroid.
If any reassignment took place, go to STEP 4, otherwise go to FIN.




KMeans

```
1 from sklearn.cluster import KMeans  
2  
3 kmeans = KMeans(n_clusters=3)  
4 kmeans.fit(X)  
5 kmeans.predict(X)
```

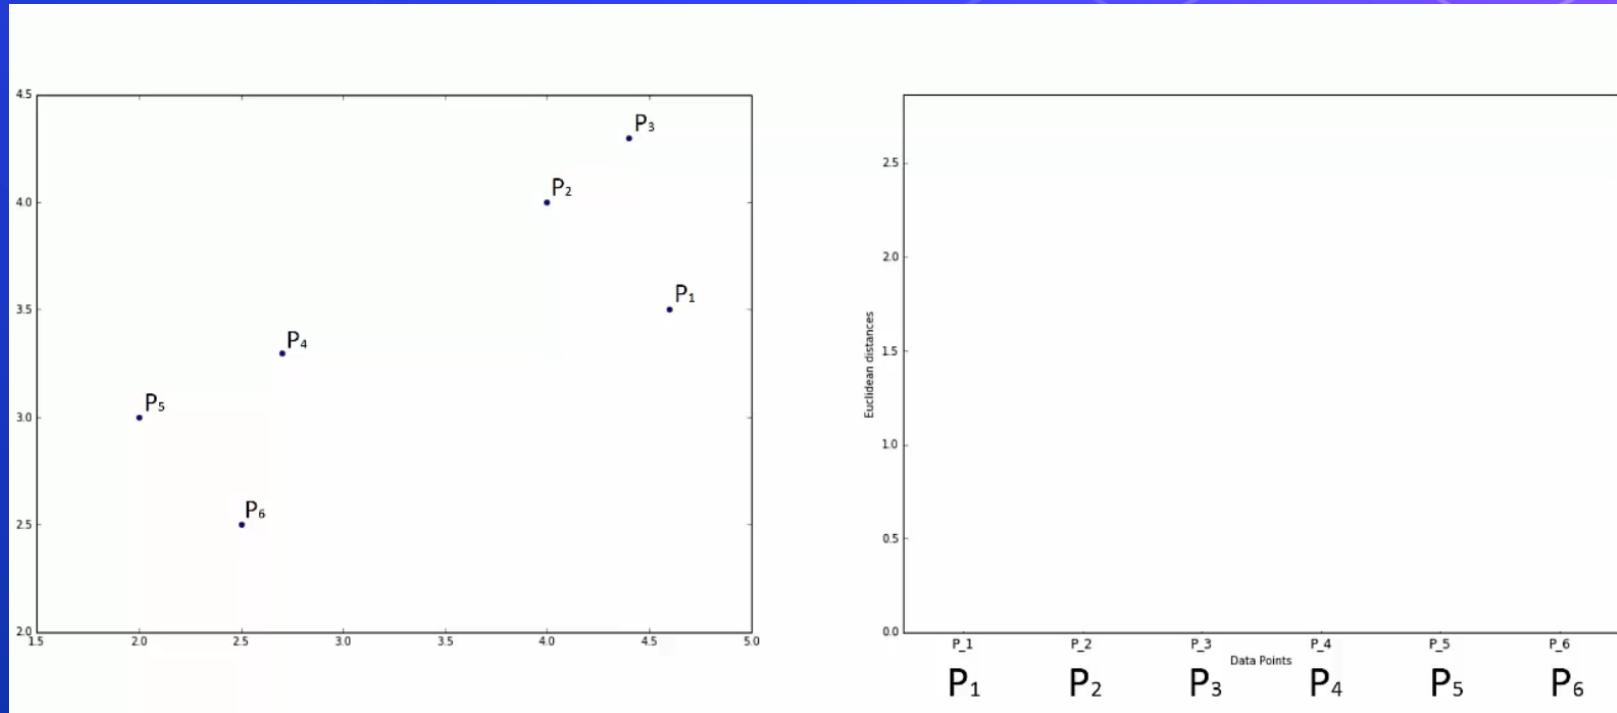
Clustering

○ KMeans


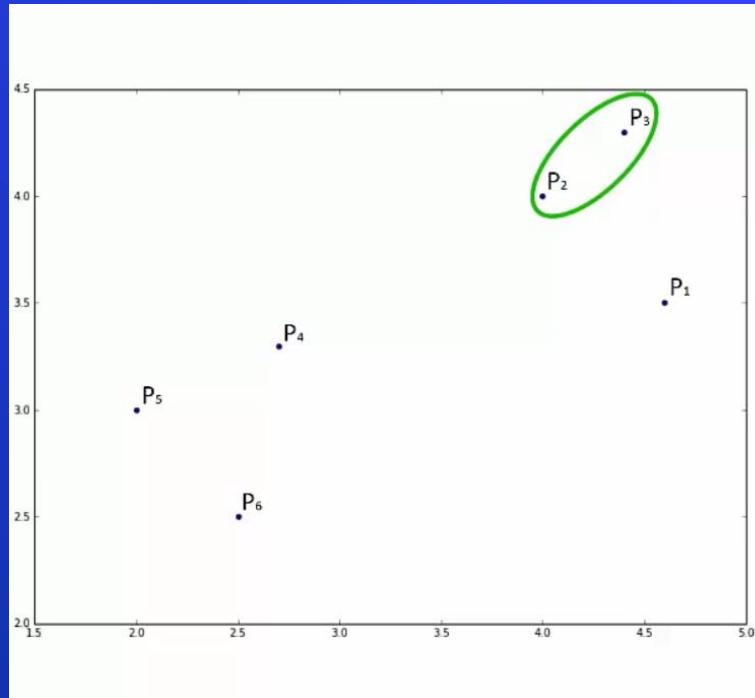
○ Hierarchical Clustering



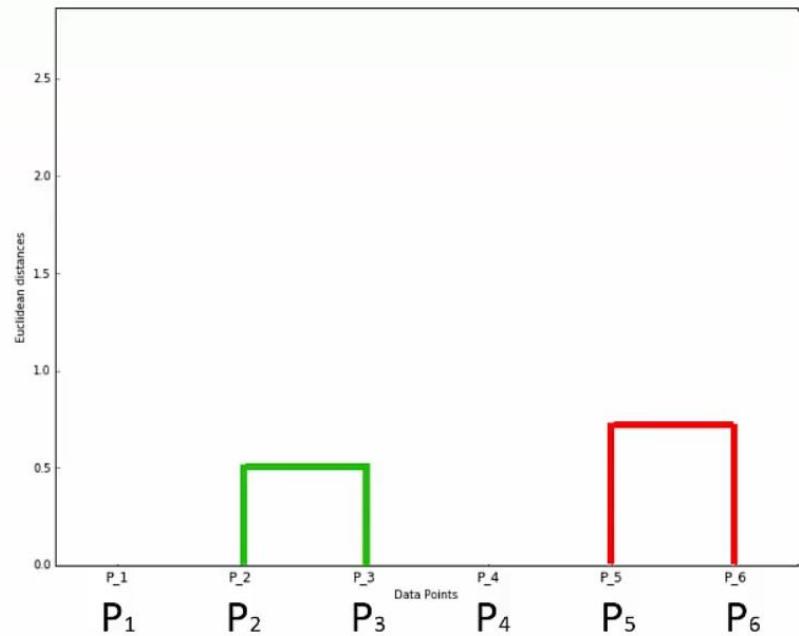
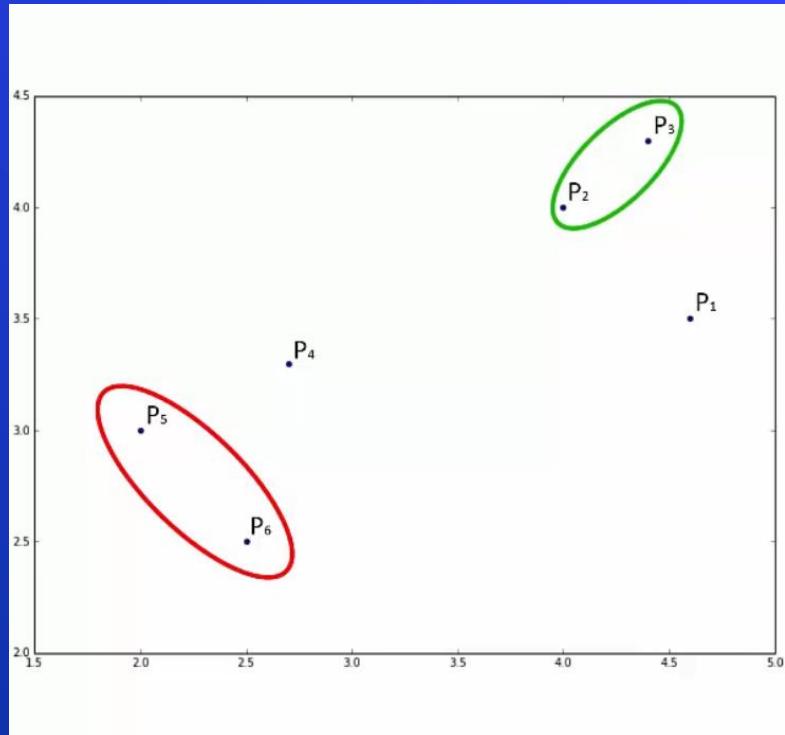
Hierarchical Clustering with dendrogram




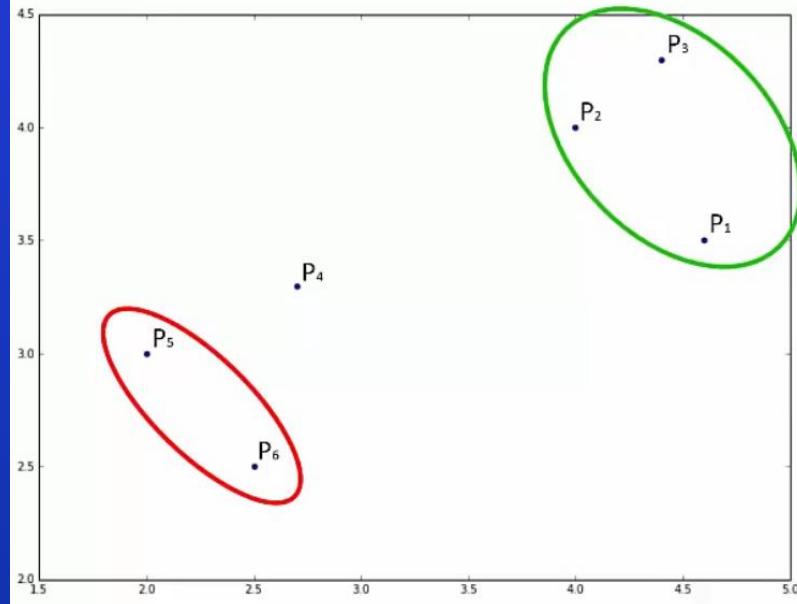
Hierarchical Clustering with dendrogram




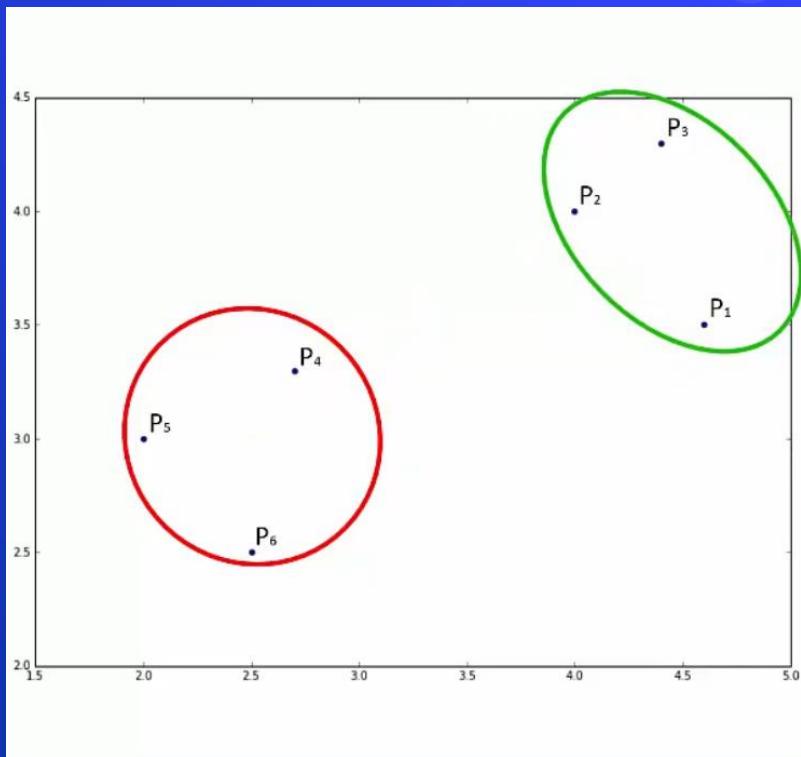
Hierarchical Clustering with dendrogram




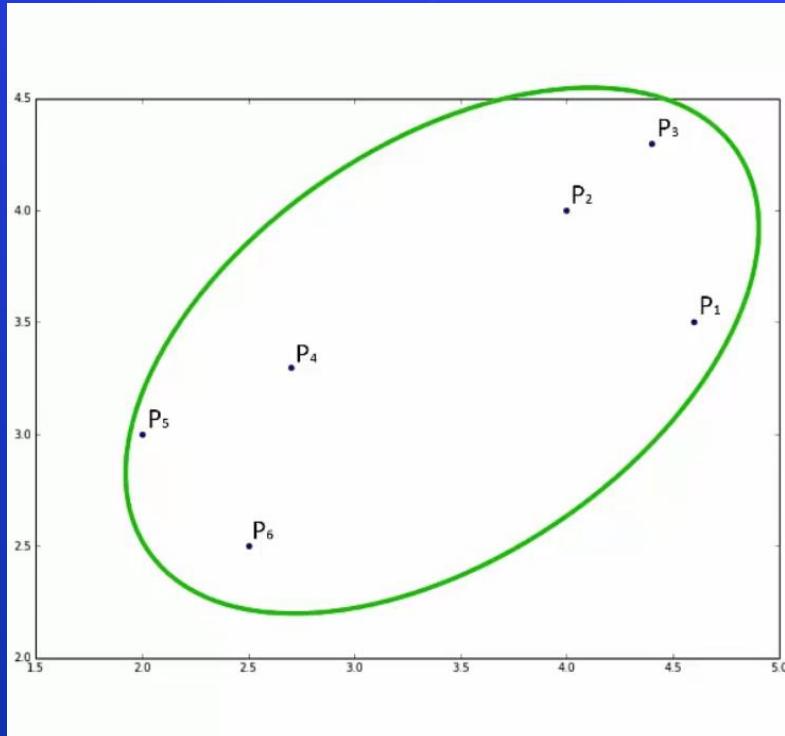
Hierarchical Clustering with dendrogram



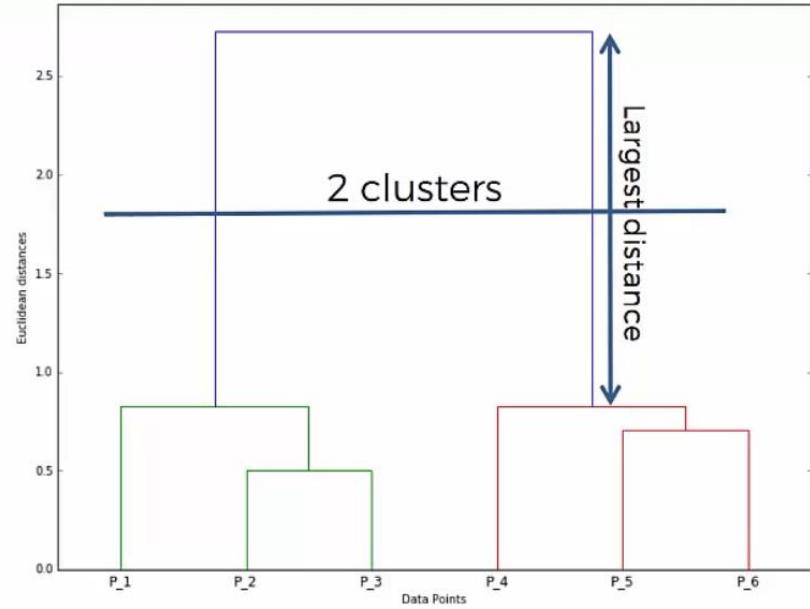

Hierarchical Clustering with dendrogram



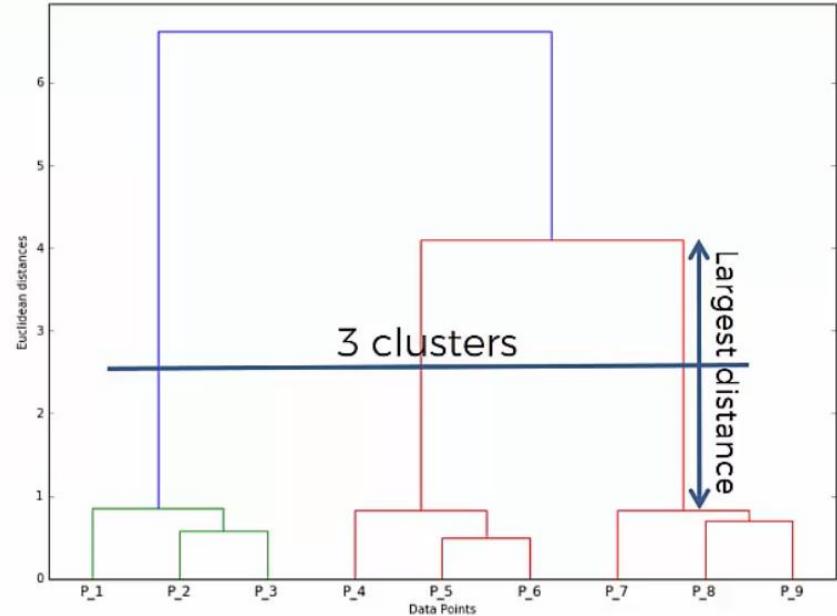
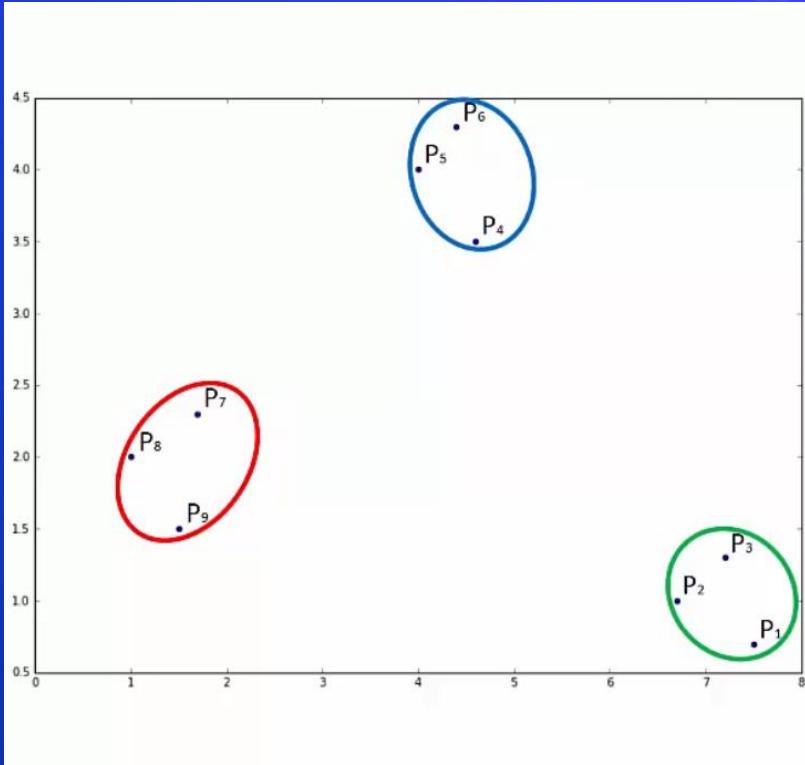
Hierarchical Clustering with dendrogram



Hierarchical Clustering with dendrogram



Hierarchical Clustering with dendrogram



Hierarchical Clustering with dendrogram

```
1 import scipy.cluster.hierarchy as sch
2 from sklearn.cluster import AgglomerativeClustering
3
4 # visualize dendrogram
5 dendrogram = sch.dendrogram(sch.linkage(x, method='ward'))
6
7 # train model
8 model = AgglomerativeClustering(n_clusters=3)
9 y_labels = model.fit_predict(x)
```

Agenda


- What is Machine Learning
- Supervised Learning
 - Regression
 - Classification
- Unsupervised Learning
 - Clustering
 - Dimension Reduction
- Model Selection & Evaluation
 - Cross Validation
 - Hyperparameter Tuning
- Intro to Deep Learning



Dimension Reduction


is the transformation of data from a high-dimensional space into a low-dimensional space so that the low-dimensional representation retains some meaningful properties of the original data

For example,
it can be used as data compression to reduce size of data for
easy processing.



Dimension Reduction




○ PCA



PCA

Data Compression

Reduce data from 3D to 2D



Andrew Ng

PCA



```
1 from sklearn.decomposition import PCA  
2  
3 X = # feature vector  
4  
5 pca = PCA(0.9)  
6 X = pca.fit_transform(X)  
7 pca.explained_variance_ratio_
```

PCA

PRINCIPAL COMPONENT ANALYSIS

PCA projects the features onto the principal components. The motivation is to reduce the features dimensionality while only losing a small amount of information.

The diagram illustrates the process of projecting data points onto principal components. The first principal component is a straight line that captures the maximum variance of the data, while the second principal component is a curved line that captures the remaining variance.

https://en.wikipedia.org/wiki/Principal_component_analysis

Agenda

- What is Machine Learning
- Supervised Learning
 - Regression
 - Classification
- Unsupervised Learning
 - Clustering
 - Dimension Reduction
- Model Selection & Evaluation
 - Cross Validation
 - Hyperparameter Tuning
- Intro to Deep Learning




Cross Validation with K-Fold

Why we use Cross Validation

- Cross Validation uses **Stratified** technique to make sure the distribution of the categorical target is the same in every fold.
- More “efficient” use of data as every observation is used for both training and testing.
- Reduce overfitting and the model achieve the generalization.

Cross Validation with K-Fold



Cross Validation with K-Fold

```
1 from sklearn.svm import SVC
2 from sklearn.model_selection import cross_validate
3
4 svc = SVC()
5 cv_results_svc = cross_validate(svc, X, Y, cv=10, return_train_score=True)
6
7 cv_results_svc['test_score'].mean()
8 """
9 0.8036085674097743
10 """
```

Agenda

- What is Machine Learning
- Supervised Learning
 - Regression
 - Classification
- Unsupervised Learning
 - Clustering
 - Dimension Reduction
- **Model Selection & Evaluation**
 - Cross Validation
 - Hyperparameter Tuning
- Intro to Deep Learning



Hyperparameter Tuning

Grid Search

Randomized Search

```
● ● ●  
1 from sklearn.model_selection import GridSearchCV  
2  
3  
4 params = [  
5     {'C':[1, 10, 100], 'kernel':['linear', 'sigmoid', 'poly']},
6     {'C':[1, 10, 100], 'kernel':['rbf'], 'gamma':[0.5, 0.6, 0.7, 0.1, 0.01, 0.001]}
7 ]  
8  
9 grid_search = GridSearchCV(estimator=model,
10                             param_grid=params,
11                             scoring='accuracy',
12                             cv=10)
13 grid_search.fit(x_train, y_train)
14  
15  
16 print(grid_search.best_params_)
17 print(grid_search.best_params_)
```

Hyperparameter Tuning

Grid Search

Randomized Search


```
● ● ●  
1 from sklearn.model_selection import RandomizedSearchCV  
2  
3  
4 params = [  
5     {'C':[1, 10, 100], 'kernel':['linear', 'sigmoid', 'poly']},
6     {'C':[1, 10, 100], 'kernel':['rbf'], 'gamma':[0.5, 0.6, 0.7, 0.1, 0.01, 0.001]}
7 ]  
8  
9 randomized_search = RandomizedSearchCV(estimator=model,
10                                         param_grid=params,
11                                         scoring='accuracy',
12                                         cv=10)
13 randomized_search.fit(x_train, y_train)
14
15
16 print(randomized_search.best_params_)
17 print(randomized_search.best_params_)
```

Agenda


- What is Machine Learning
- Supervised Learning
 - Regression
 - Classification
- Unsupervised Learning
 - Clustering
 - Dimension Reduction
- Model Selection & Evaluation
 - Cross Validation
 - Hyperparameter Tuning
- Intro to Deep Learning




What is Deep Learning



What is Deep Learning




Traditional machine learning uses hand-crafted features, which is tedious and costly to develop.





Deep learning learns hierarchical representation from the data itself, and scales with more data.

What is Deep Learning



What is Deep Learning




What is Deep Learning

Why is Deep Learning Hot Now?

Big Data Availability



350 millions
images uploaded
per day




2.5 Petabytes of
customer data
hourly



300 hours of video
uploaded every
minute

New ML Techniques




GPU Acceleration



Deep Learning Applications

DL Today: Speech-to-Text



[Baidu 2014]



Deep Learning Applications

DL Today: Vision



[Krizhevsky 2012]



[Ciresan et al. 2013]



[Faster R-CNN - Ren 2015]



[NVIDIA dev blog]




Deep Learning Applications

DL Today: Vision



[Stanford 2017]



[Nvidia Dev Blog 2017]




Figure 1. Illumination and Pose invariance.

[FaceNet - Google 2015]




[Facial landmark detection CUHK 2014]



Deep Learning Applications

DL Today: NLP



Deep Learning Applications

DL Today: NLP



Salit Kulla
to me

11:29 AM ***

Hey, Wynton Marsalis is playing this weekend. Do you have a preference between Saturday and Sunday?

-S

I'm down for either.

Let's do Saturday.

I'm fine with whatever.

Reply

Forward

[Google Inbox Smart Reply]




[Amazon Echo / Alexa]



Deep Learning Applications

DL Today: Vision + NLP



[VQA - Mutan 2017]



"man in black shirt is playing guitar."



"construction worker in orange safety vest is working on road."



"two young girls are playing with lego toy."



"boy is doing backflip on wakeboard."

[Karpathy 2015]



Deep Learning Applications


DL Today: Image translation



[DeepDream 2015]



[Gatys 2015]



bicubic
(21.59dB/0.6423)



SRResNet
(23.44dB/0.7777)



SRGAN
(20.34dB/0.6562)

[Ledig 2016]



Deep Learning Applications

DL Today: Generative models



Sampled celebrities [Nvidia 2017]


Text description	This bird is blue with white and has a very short beak	This bird has wings that are brown and has a yellow belly	A white bird with a black crown and yellow beak	This bird is white, black, and brown in color, with a brown beak	The bird has small beak, with reddish brown crown and gray belly	This is a small, black bird with a white breast and white on the wingbars.	This bird is white black and yellow in color, with a short black beak
Stage-I images							
Stage-II images							

StackGAN v2 [Zhang 2017]



Deep Learning Applications

DL Today: Generative models



Sound generation with WaveNet [DeepMind 2017]




Deep Learning Applications


DL for AI in games



[Deepmind AlphaGo / Zero 2017]



A10: English Opening
w 20/30(0), b 8/40(2)



[Atari Games - DeepMind 2016]



[Starcraft 2 for AI research]



Questions ?!



Thanks!

>_ Live long and prosper

